



### Annis on MonetDB

Viktor Rosenfeld rosenfel@informatik.hu-berlin.de

14. January 2013

Advisors: Prof. Dr. Ulf Leser and Dr. Stefan Manegold



http://www.flickr.com/photos/karola/3623768629

- 1. What is Annis and how is it used?
- 2. Current implementation on PostgreSQL
- 3. What are Column-Stores? How can Annis benefit?
- 4. New implementation on MonetDB and evaluation

1. What is Annis and how is it used?

## What's a corpus?

any principled collection of language

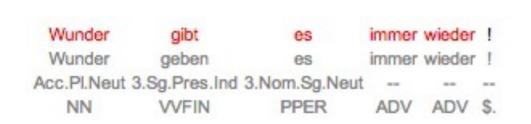


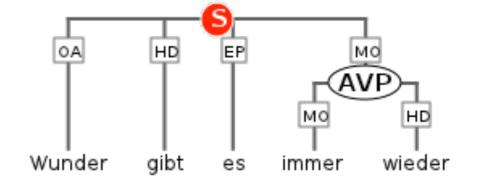






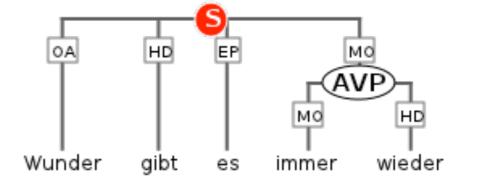
# classification and interpretation of the corpus data additional data to enrich the corpus





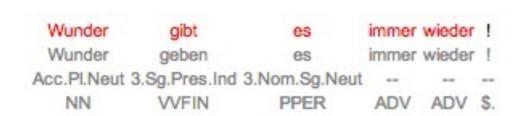
# classification and interpretation of the corpus data additional data to enrich the corpus

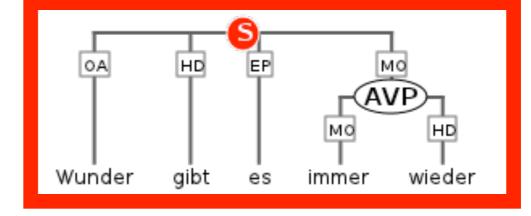




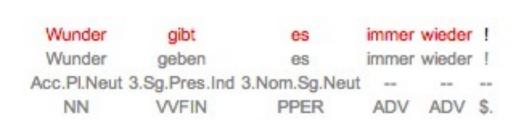
classification and interpretation of the corpus data

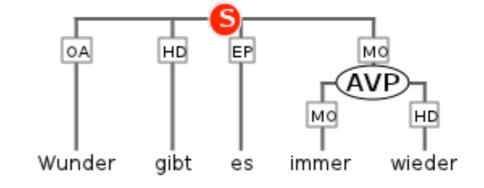
additional data to enrich the corpus





classification and interpretation of the corpus data additional data to enrich the corpus

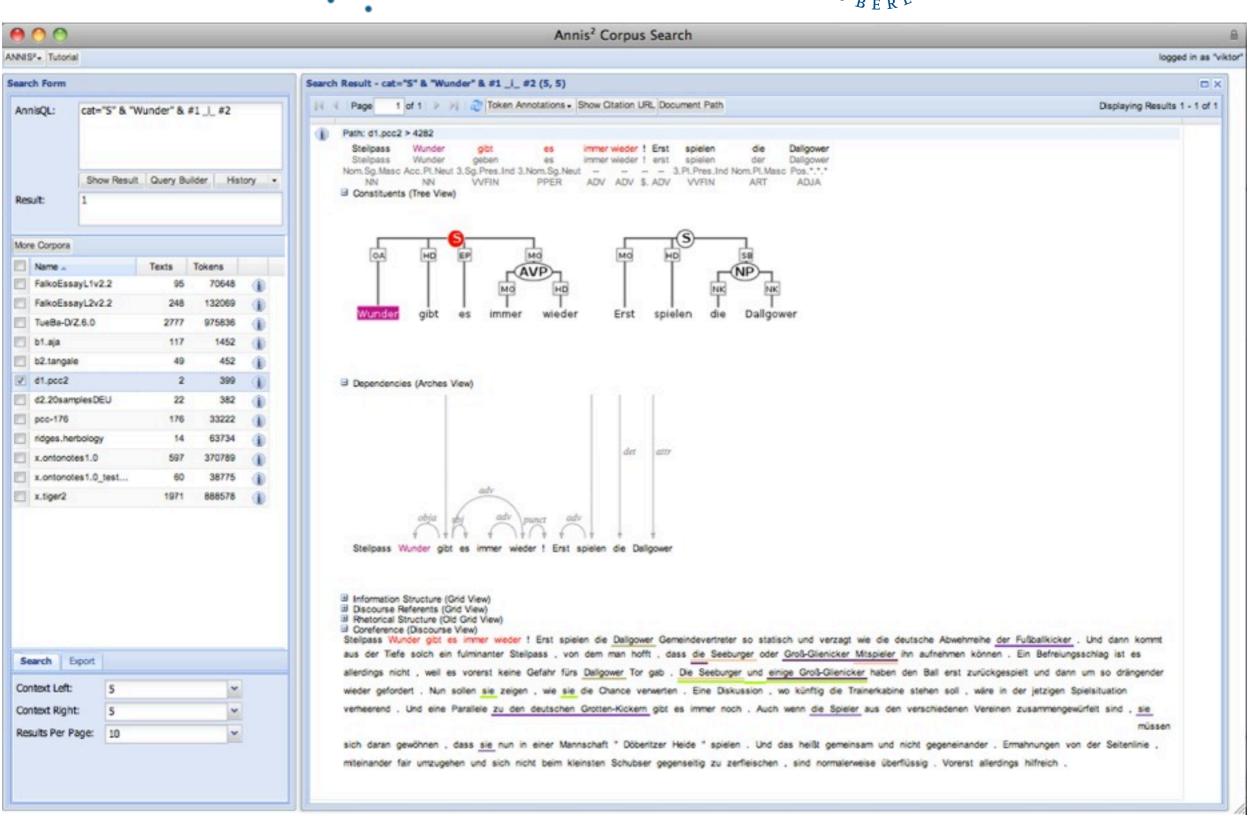






### Annis

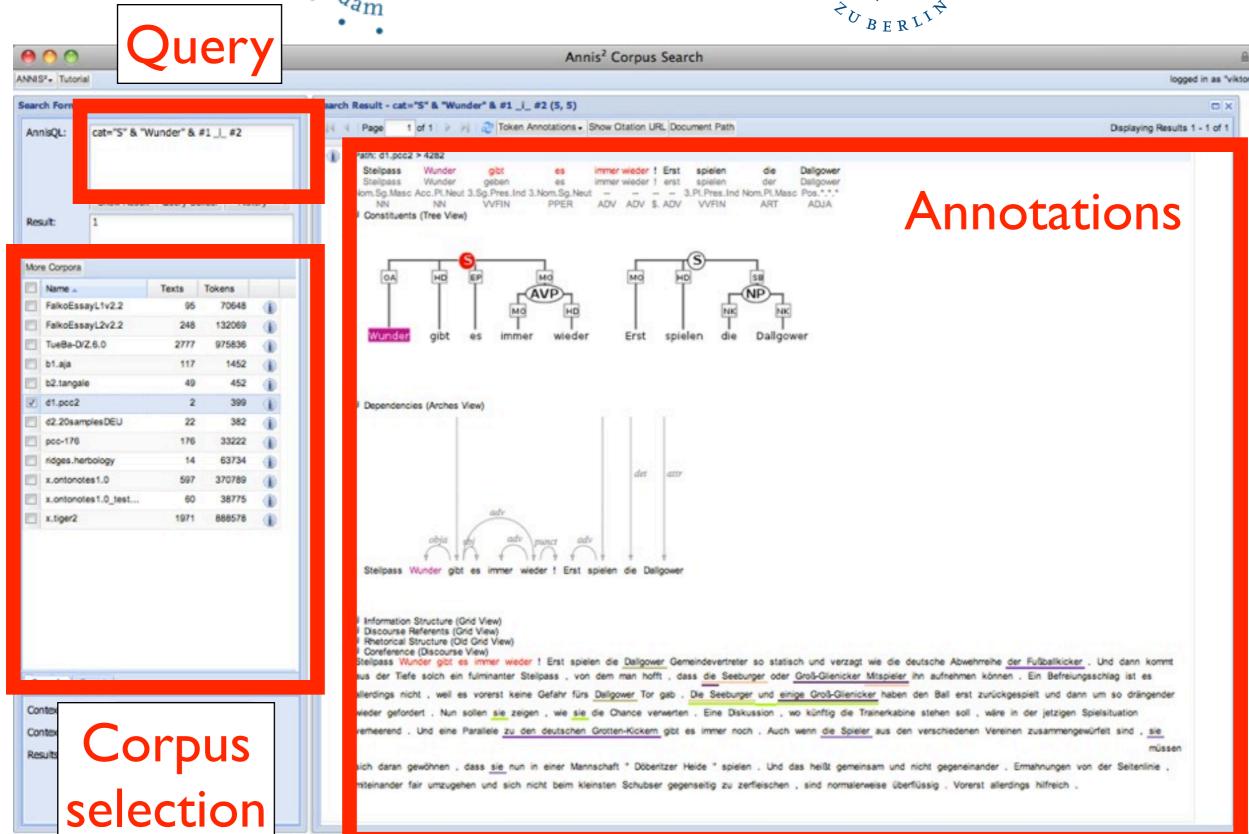






### Annis

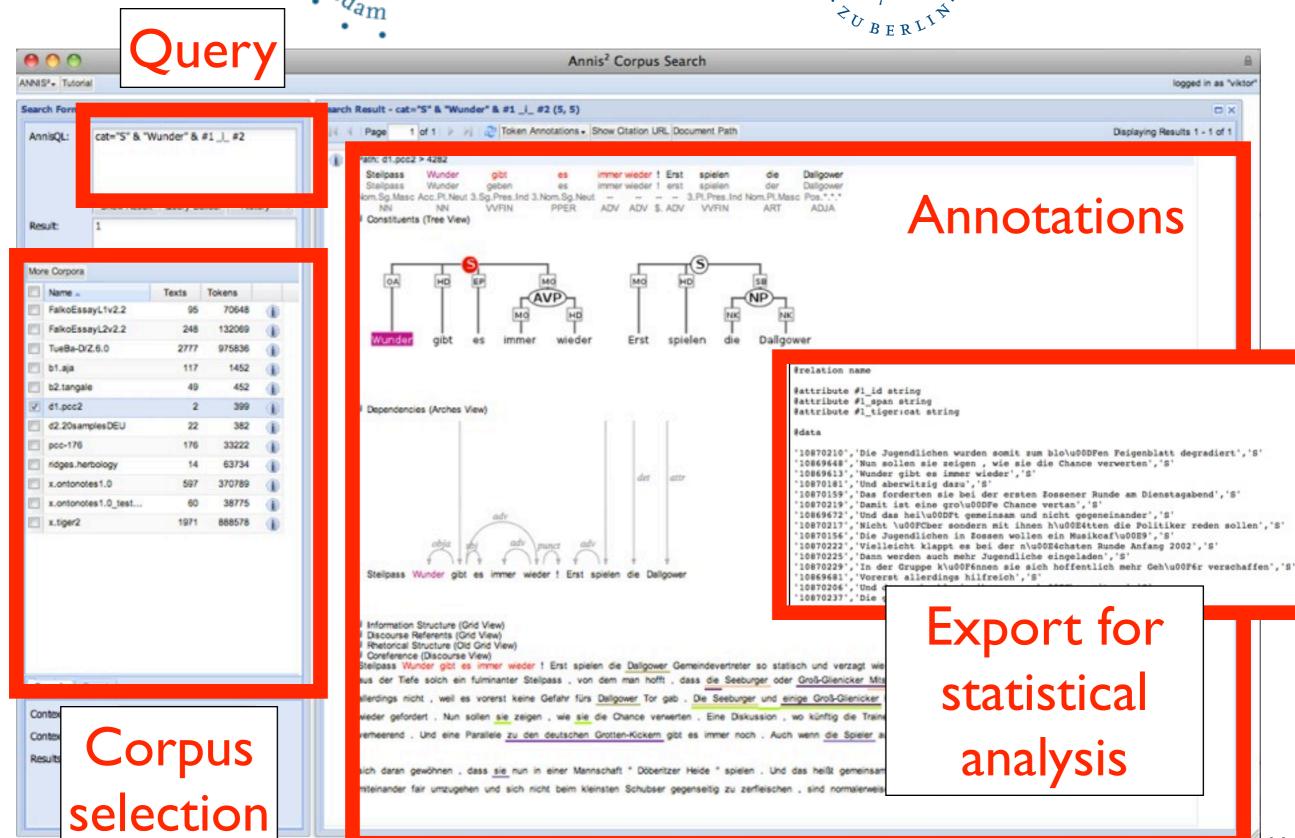






### Annis





## Annis query language

```
cat="S" & find a sentence
"Wunder" & and find the phrase "Wunder"
#1 _i_ #2 the sentence includes the phrase "Wunder"
```

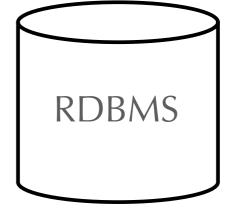
## Annis query language

```
cat="S" & find a sentence
"Wunder" & and find the phrase "Wunder"
#1 _i_ #2 the sentence includes the phrase "Wunder"
```



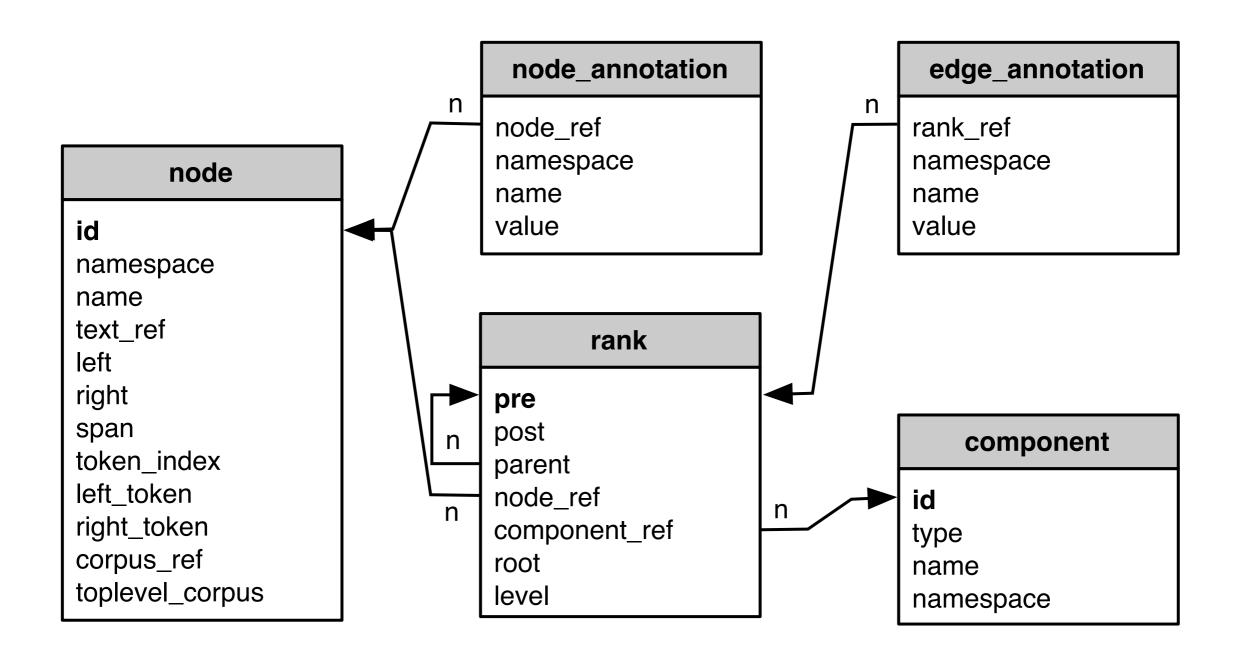
SELECT id1, id2
FROM ...
WHERE ...



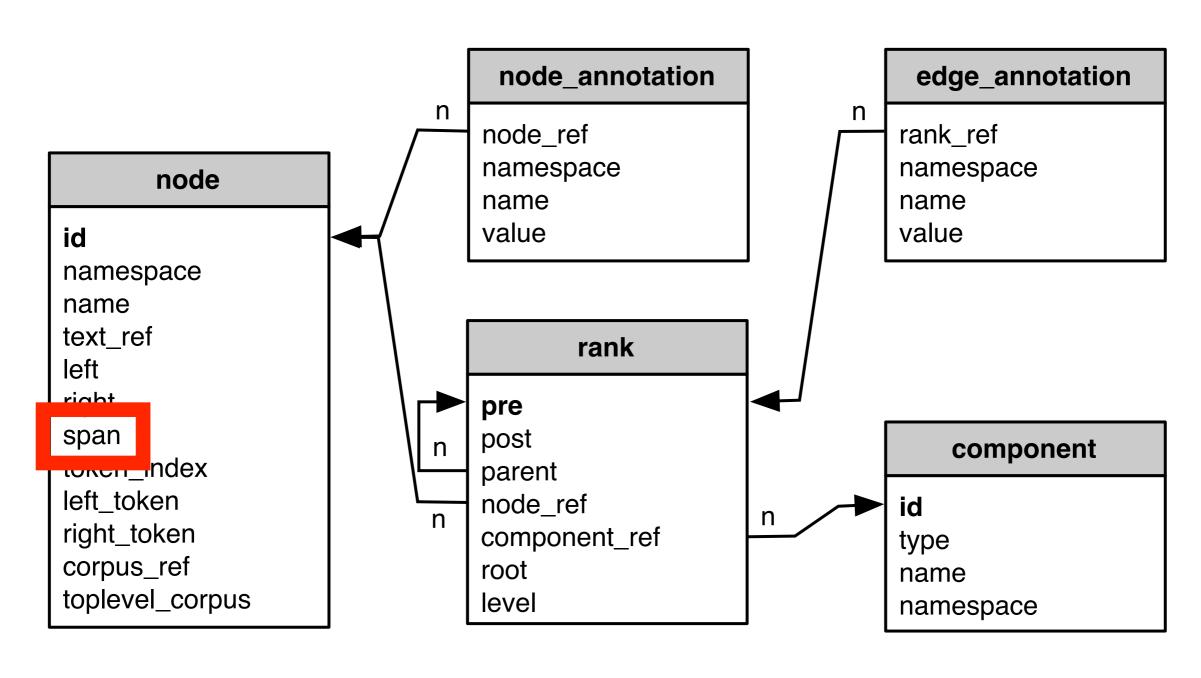


2. Current implementation on PostgreSQL

### Database schema



## Example 1: Text search

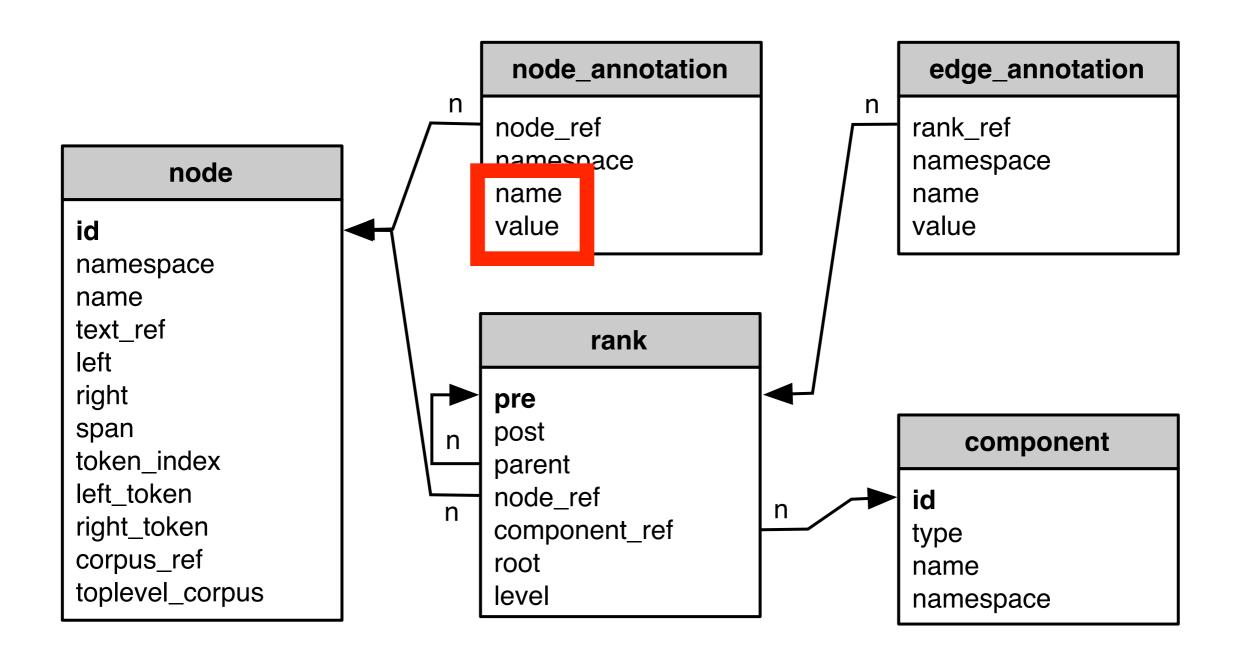


AQL: "Wunder"



SQL:nodeN.span = 'Wunder'

## Example 2: Annotation search

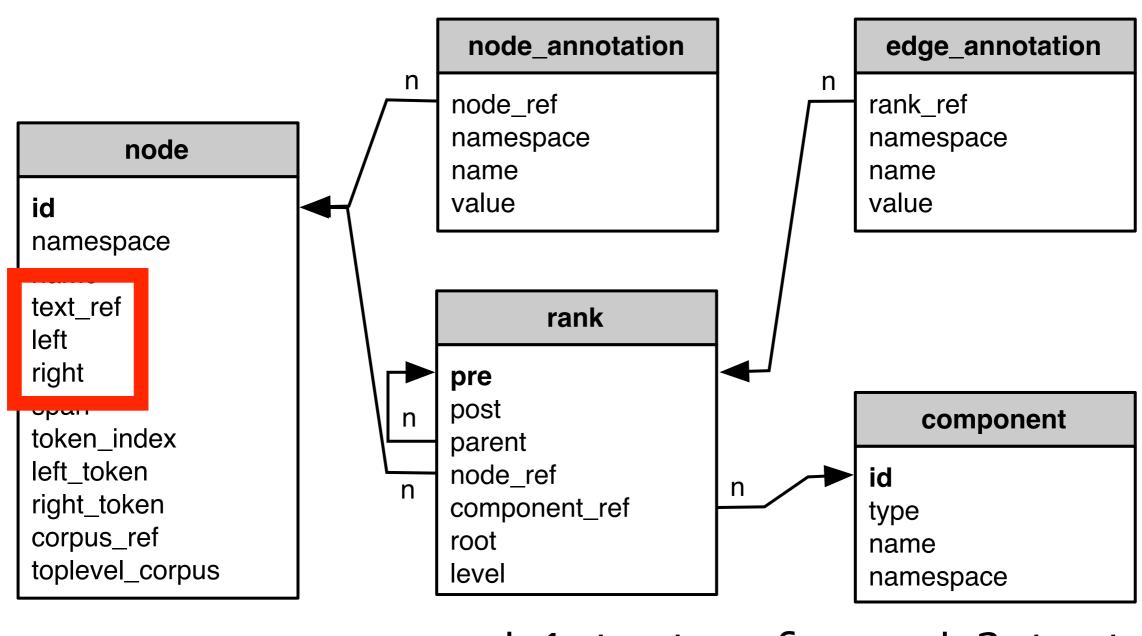






node\_annotationN.name = 'cat'
node\_annotationN.value = 'S'

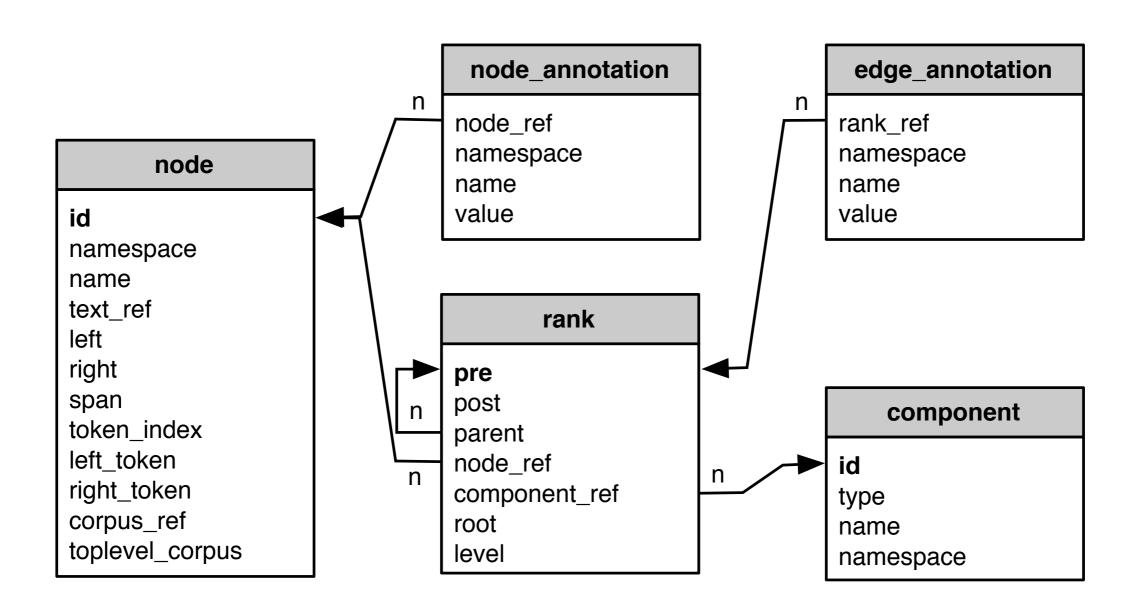
## Example 3: Inclusion operator



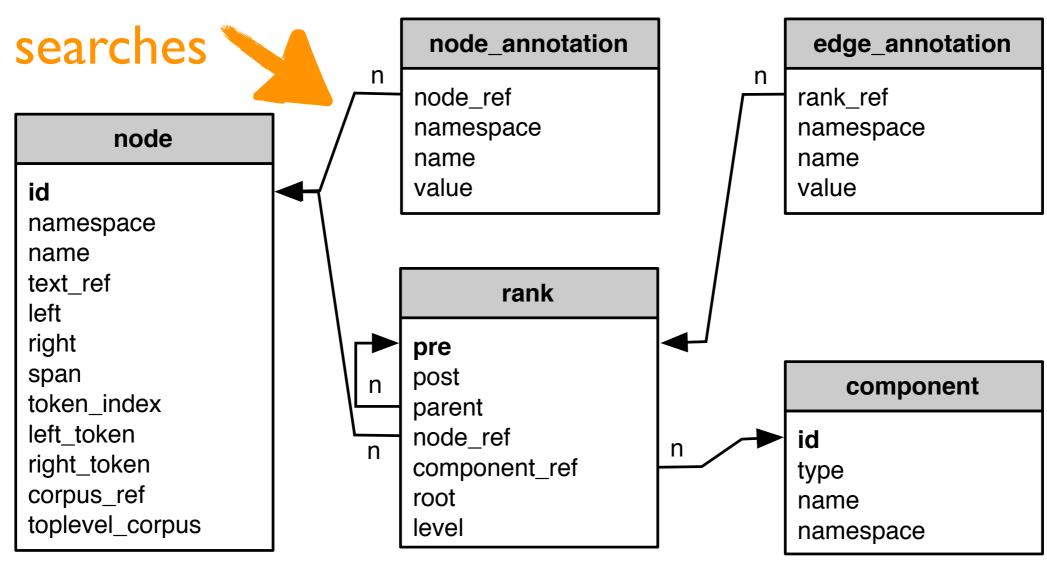
#1 \_i\_ #2



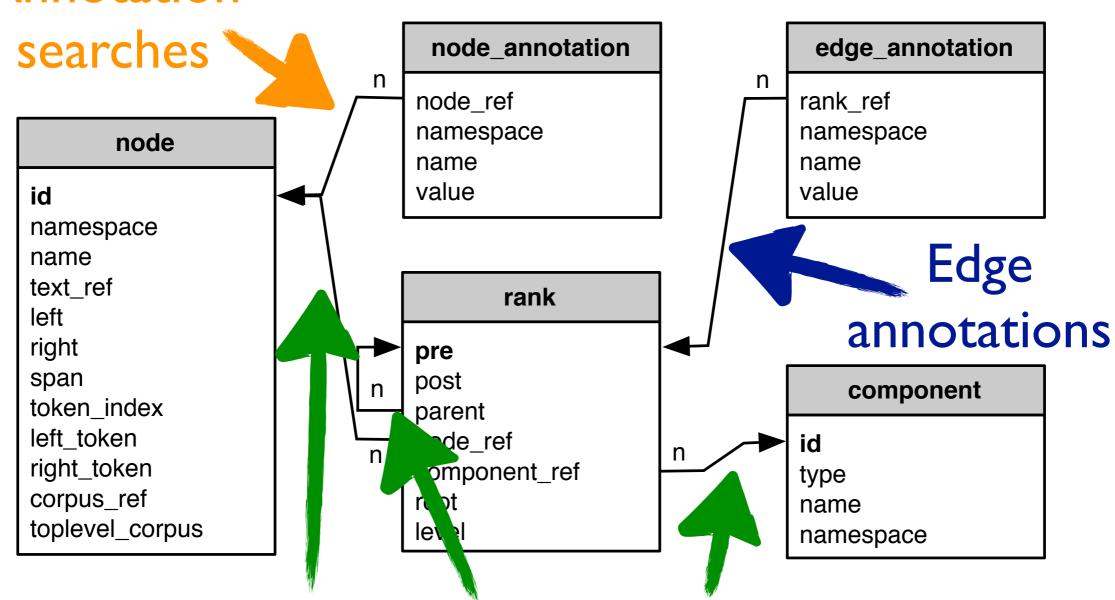
node1.text\_ref = node2.text\_ref
node1.right <= node2.right
node1.left >= node2.left



#### **Annotation**

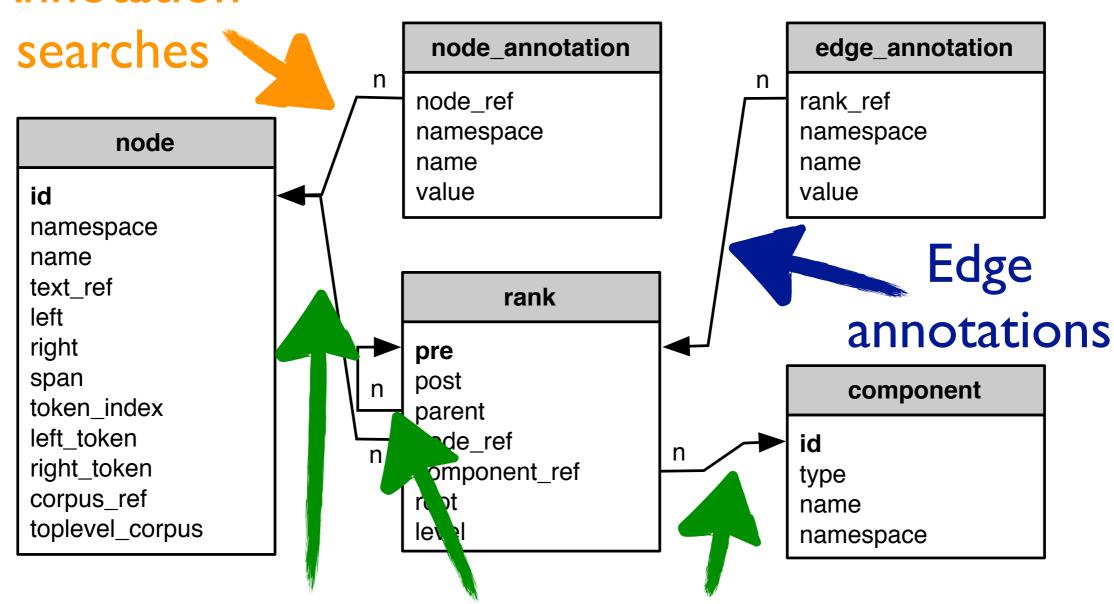


#### **Annotation**

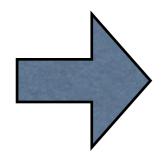


Binary relations on edges

#### **Annotation**

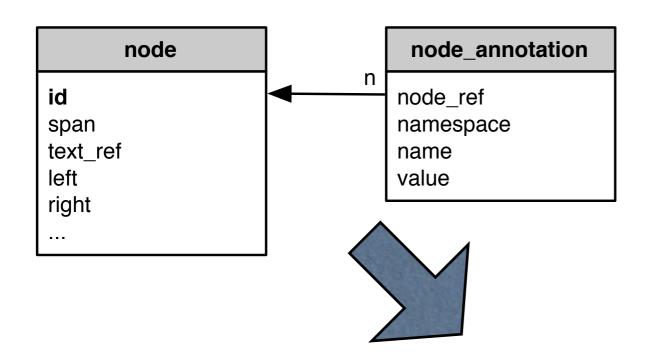


Binary relations on edges



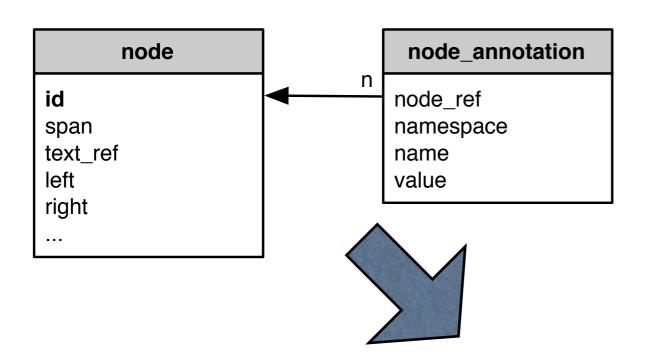
Bad performance on PostgreSQL

## Solution 1: One big table



id	span	text_ref	left	right	 na_name	na_value	•••
Ι		I		30	cat	S	
2	Wunder	I	I	5	morph	Acc.Pl.Neut	
2	Wunder	I	I	5	pos	NN	
2	Wunder	I	I	5	lemma	Wunder	
•••							

## Solution 1: One big table



id	span	text_ref	left	right	 na_name	na_value	•••
				30	cat	S	
2	Wunder	I	I	5	morph	Acc.Pl.Neut	
2	Wunder	I	I	5	pos	NN	
2	Wunder	Ī	Ī	5	lemma	Wunder	
•••							

Pro: Fewer joins

Contra: Increased redundancy, less extensible

### Solution 2: Combined indexes

id	span	text_ref	left	right	 na_name	na_value	•••
		I		30	cat	S	
2	Wunder	I		5	morph	Acc.Pl.Neut	
2	Wunder	l		5	pos	NN	
2	Wunder	I		5	lemma	Wunder	
•••							

#### One index over 4 columns

#### Find nodes

- spanning a certain word,
- in a certain text,
- at a certain position.

Pro: Potentially very fast Contra: Uses lots of disk space

## Disk usage in PostgreSQL

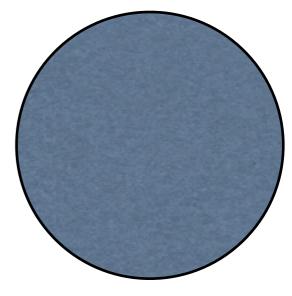
#### TIGER Treebank 2.1

ca. 50.000 sentences, 900.000 tokens, 3 million annotations, 1 million edges





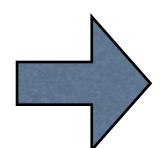




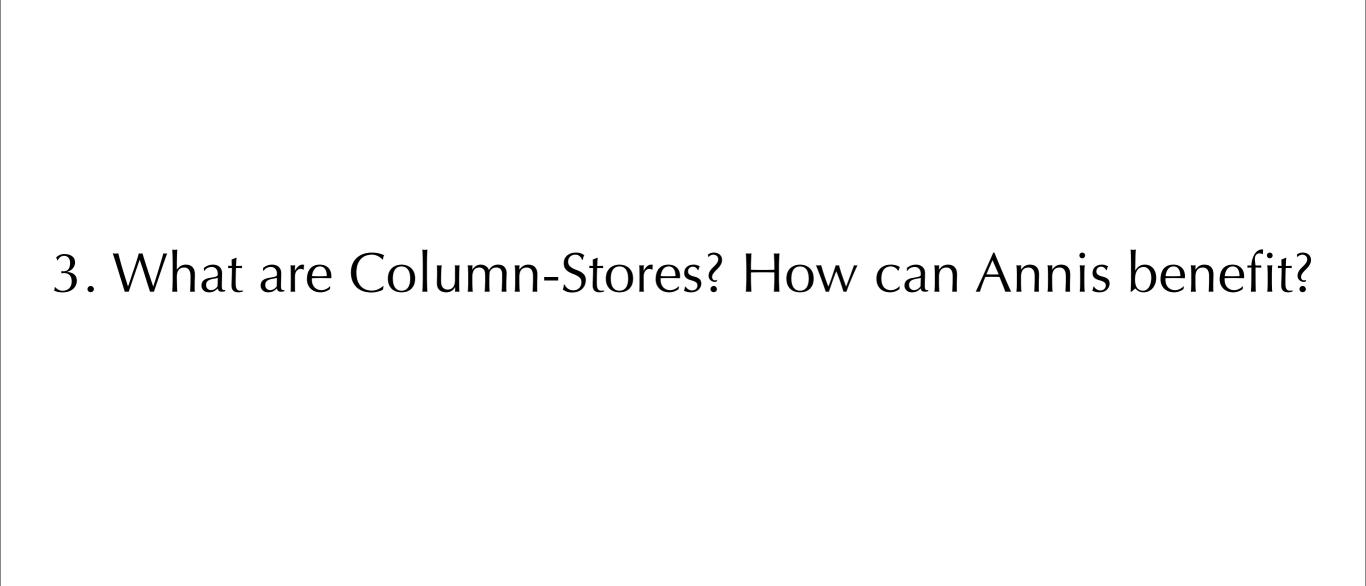
280 MB Data files 525 MB Normalized (many tables)

I.2 GB Materialized (one table)

7.7 GB
Materialized
+ Indexes



Increase by factor 15 (or almost 30)



### What's a Column-Store?

conceptual model

	node_ref	name	value
I	123	pos	VVINF
2	123	lemma	essen
3	456	pos	NN

table

storage model

### What's a Column-Store?

conceptual model

	node_ref	name	value
I	123	pos	VVINF
2	123	lemma	essen
3	456	pos	NN

table



storage model

I	123	pos	VVINF
2	123	lemma	essen
3	456	pos	NN

rows

### What's a Column-Store?

conceptual model

	node_ref	name	value
ı	123	pos	VVINF
2	123	lemma	essen
3	456	pos	NN

table





storage model

I	123	pos	VVINF
2	123	lemma	essen
3	456	pos	NN

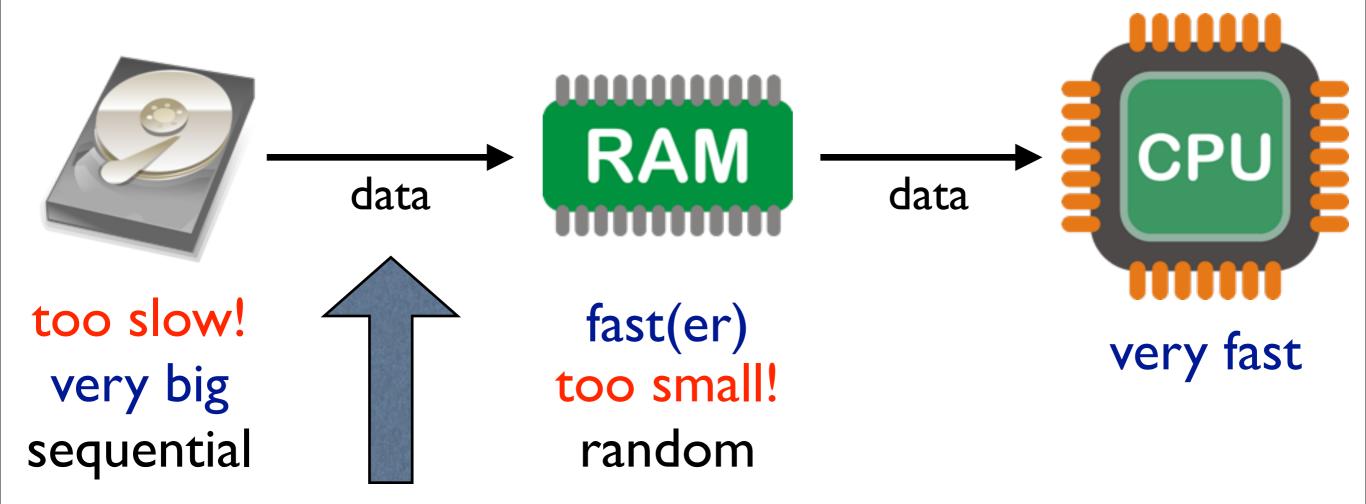
rows

node_ref	name
123	pos
123	lemma
456	pos

value
VVINF
essen
NN

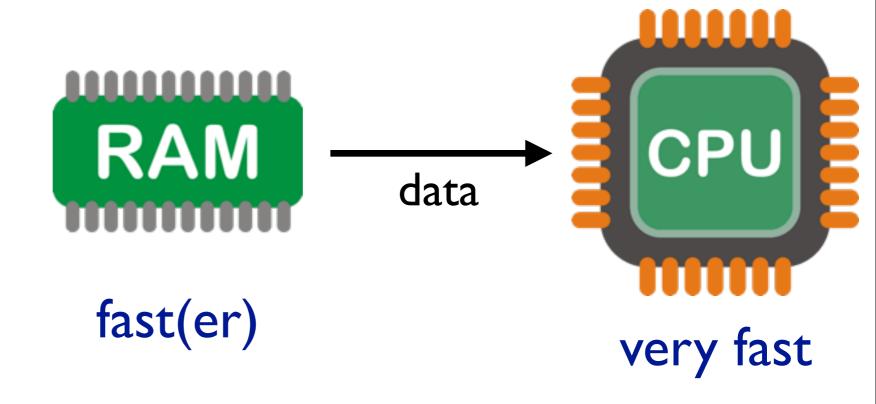
columns

## Why Column-Stores? – Why Databases?

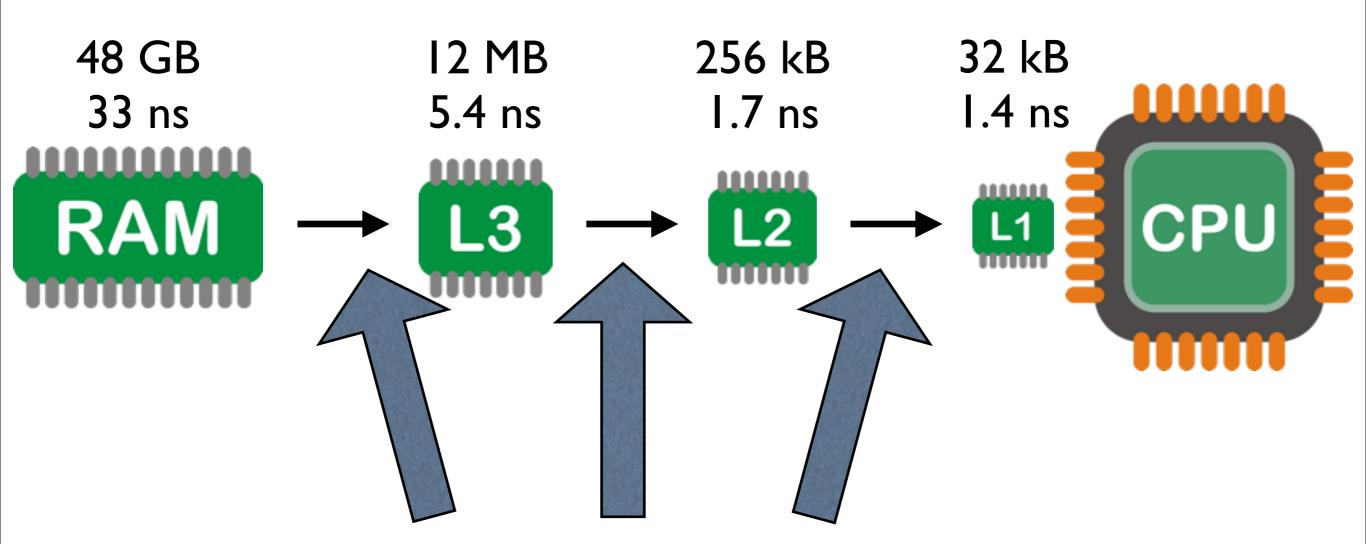


job of the database (traditionally)

### Why Column-Stores? – Why Databases?



### Caches between RAM and CPU



job of the database on a modern system (among others)

## Cache usage of row layout

query: compare name attribute with value 'lemma'

data file: I 123 pos VVINF 2 123 lemma essen 3 456 pos

L1 cache:

NN

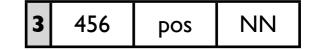
## Cache usage of row layout

query: compare name attribute with value 'lemma'

data file:





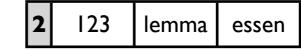


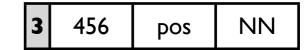
I. load first row

query: compare name attribute with value 'lemma'

data file:





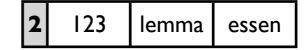


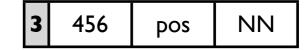
- I. load first row
- 2. locate name attribute

query: compare name attribute with value 'lemma'

data file:





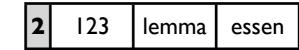


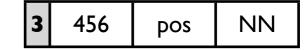
- I. load first row
- 2. locate name attribute
- 3. test name attribute

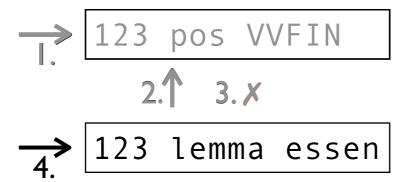
query: compare name attribute with value 'lemma'

data file:







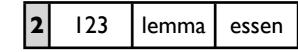


- I. load first row
- 2. locate name attribute
- 3. test name attribute
- 4. load second row

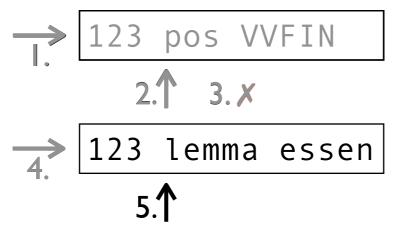
query: compare name attribute with value 'lemma'

data file:







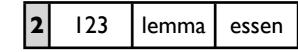


- I. load first row
- 2. locate name attribute
- 3. test name attribute
- 4. load second row
- 5. locate name attribute

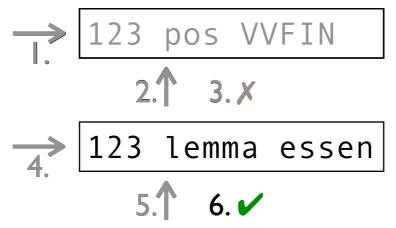
query: compare name attribute with value 'lemma'

data file:







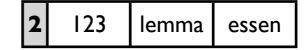


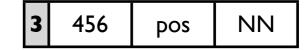
- I. load first row
- 2. locate name attribute
- 3. test name attribute
- 4. load second row
- 5. locate name attribute
- 6. test name attribute

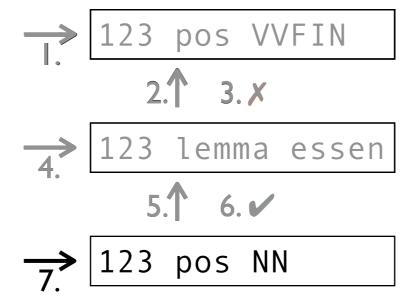
query: compare name attribute with value 'lemma'

data file:







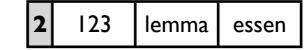


- I. load first row
- 2. locate name attribute
- 3. test name attribute
- 4. load second row
- 5. locate name attribute
- 6. test name attribute
- 7. load third row

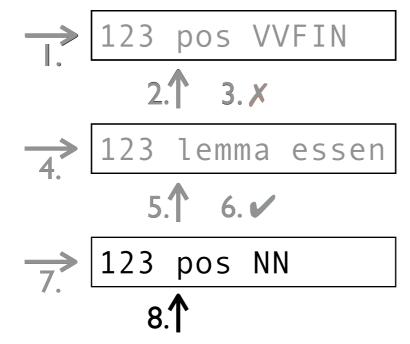
query: compare name attribute with value 'lemma'

data file:





	3	456	pos	Z
--	---	-----	-----	---



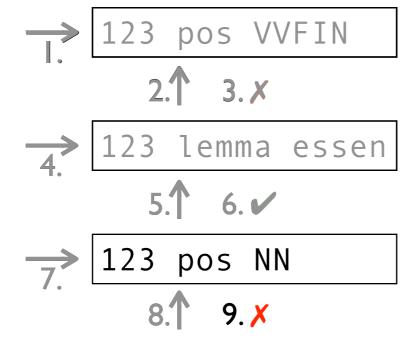
- I. load first row
- 2. locate name attribute
- 3. test name attribute
- 4. load second row
- 5. locate name attribute
- 6. test name attribute
- 7. load third row
- 8. locate name attribute

query: compare name attribute with value 'lemma'

data file:



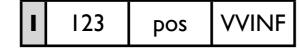
2 123 lemma essen

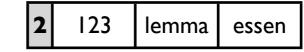


- I. load first row
- 2. locate name attribute
- 3. test name attribute
- 4. load second row
- 5. locate name attribute
- 6. test name attribute
- 7. load third row
- 8. locate name attribute
- 9. test name attribute

query: compare name attribute with value 'lemma'

data file:





	3	456	pos	ZZ
--	---	-----	-----	----

- I. load first row
- 2. locate name attribute
- 3. test name attribute
- 4. load second row
- 5. locate name attribute
- 6. test name attribute
- 7. load third row
- 8. locate name attribute
- 9. test name attribute

query: compare name attribute with value 'lemma'

data file:

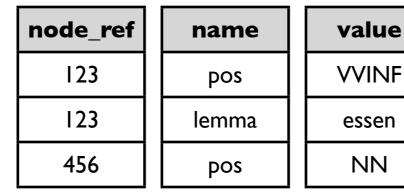
node_ref
123
123
456

name
pos
lemma
pos

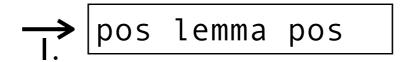
value
VVINF
essen
NN

query: compare name attribute with value 'lemma'

data file:



L1 cache:



I. load name column

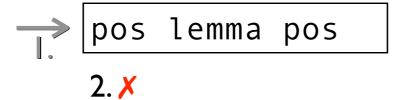
query: compare name attribute with value 'lemma'

data file:

node_ref	
123	
123	
456	

name	
pos	
lemma	
pos	

value	
VVINF	
essen	
NN	



- I. load name column
- 2. test first name attribute

query: compare name attribute with value 'lemma'

data file:

node_ref
123
123
456

name
pos
lemma
pos

value
VVINF
essen
NN



- I. load name column
- 2. test first name attribute
- 3. test second name attribute

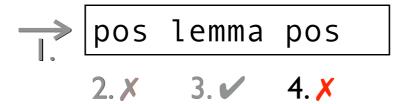
query: compare name attribute with value 'lemma'

data file:

node_ref
123
123
456

name
pos
lemma
pos

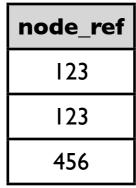
value
VVINF
essen
NN



- I. load name column
- 2. test first name attribute
- 3. test second name attribute
- 4. test third name attribute

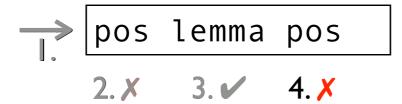
query: compare name attribute with value 'lemma'

data file:



name
pos
lemma
pos

value
VVINF
essen
NN



- I. load name column
- 2. test first name attribute
- 3. test second name attribute
- 4. test third name attribute

#### Column operations in Annis

#### Search terms

can be indexed

"Wunder"

#### Regular expressions

- can often be indexed
- but not always

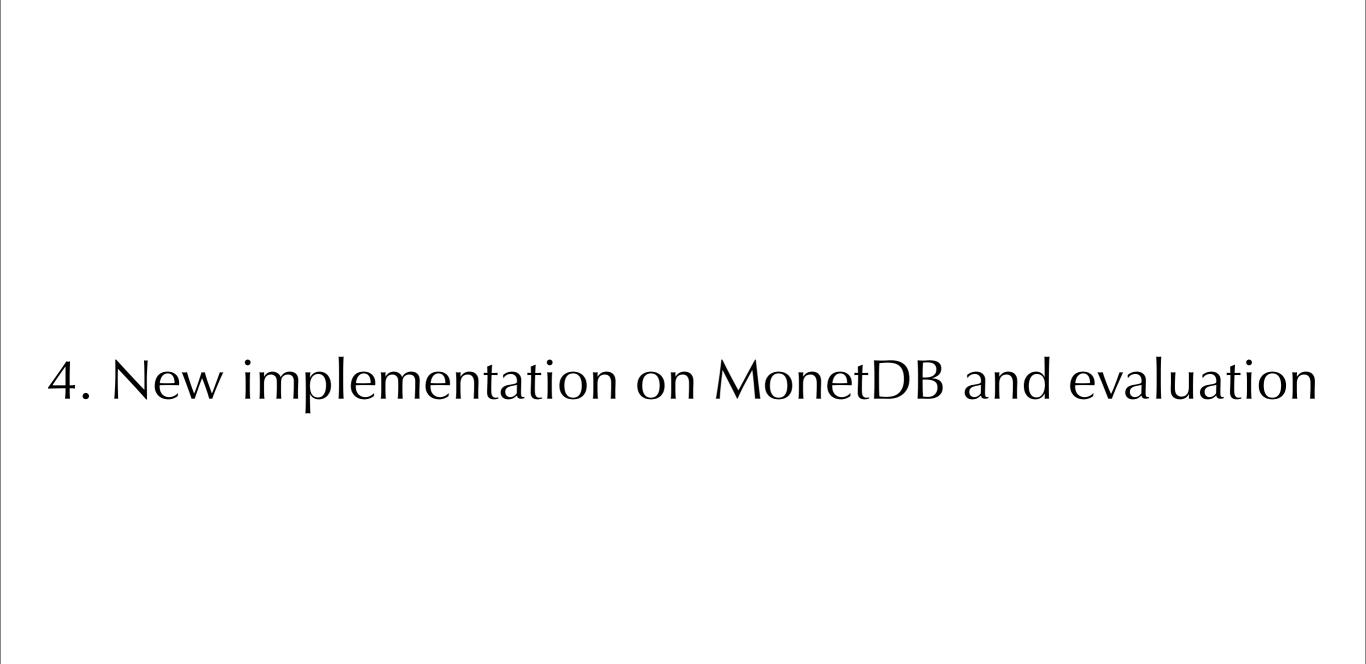
morph=/.\*\.Pl\.Neut/

#### Binary operations

- can be indexed
- need many indexes
- slow if there are many index lookups

	•	
_=_	_1_	>

span	text ref	left	right	
			30	
Wunder			5	
Wunder			5	
Wunder			5	
	Wunder Wunder	Wunder I Wunder I	I I Wunder I I Wunder I I	I



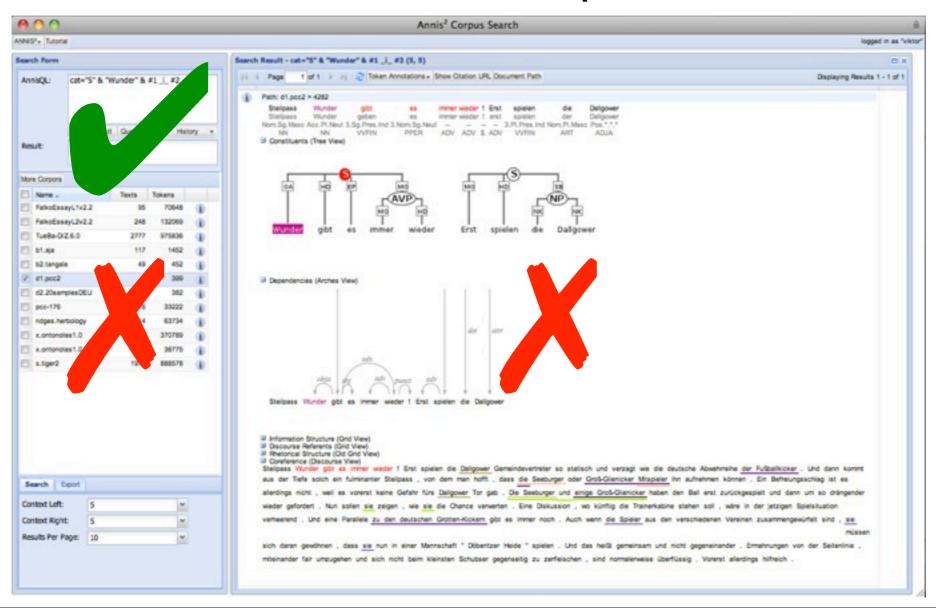
#### Prototype implementation

#### Supported

- Annis 2 Query Language
- COUNT queries

#### **Not supported**

- Annis 3 language features
- ANNOTATE, MATRIX queries
- corpus selection



#### Realistic test workload

Corpus: TIGER Treebank 2.1

#### **Queries:**

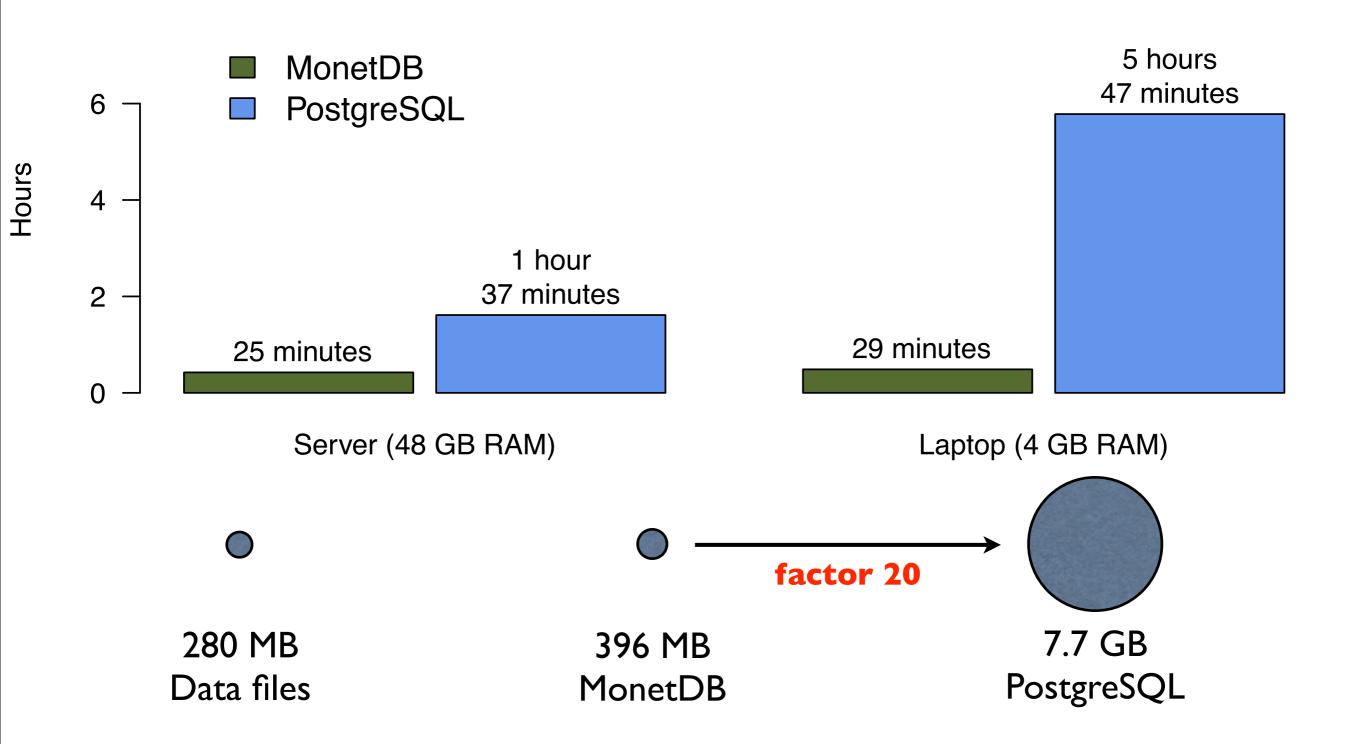
- 3 month query log of Annis instance at the SFB 632
- 337 TIGER queries (224 unique)
- up to 4 search terms
- up to 6 binary operators

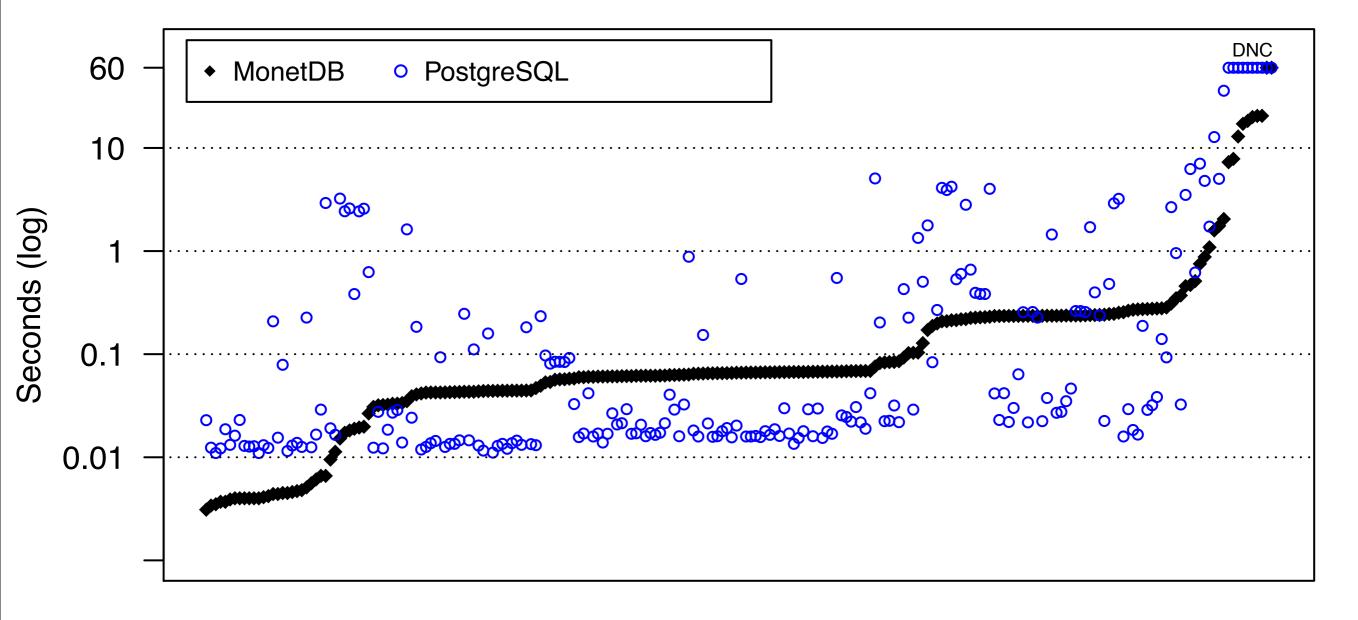
#### **Random workload:**

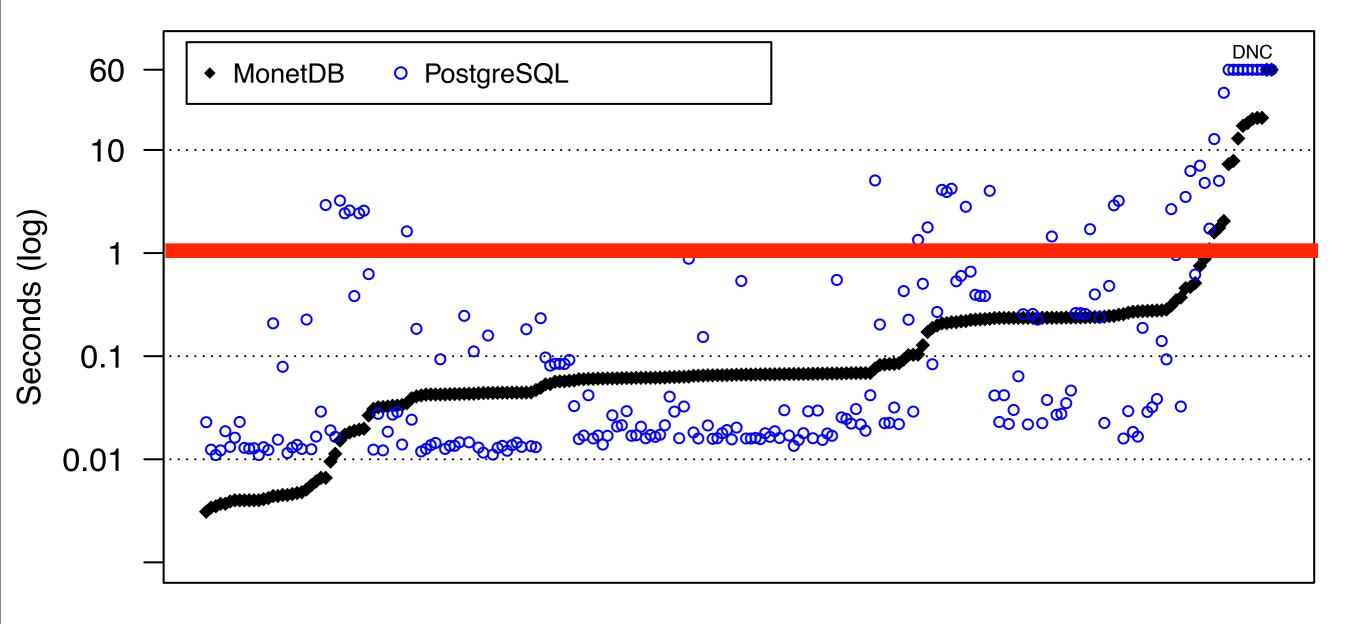
- 10000 queries
- original distribution
- excluded PostgreSQL timeout

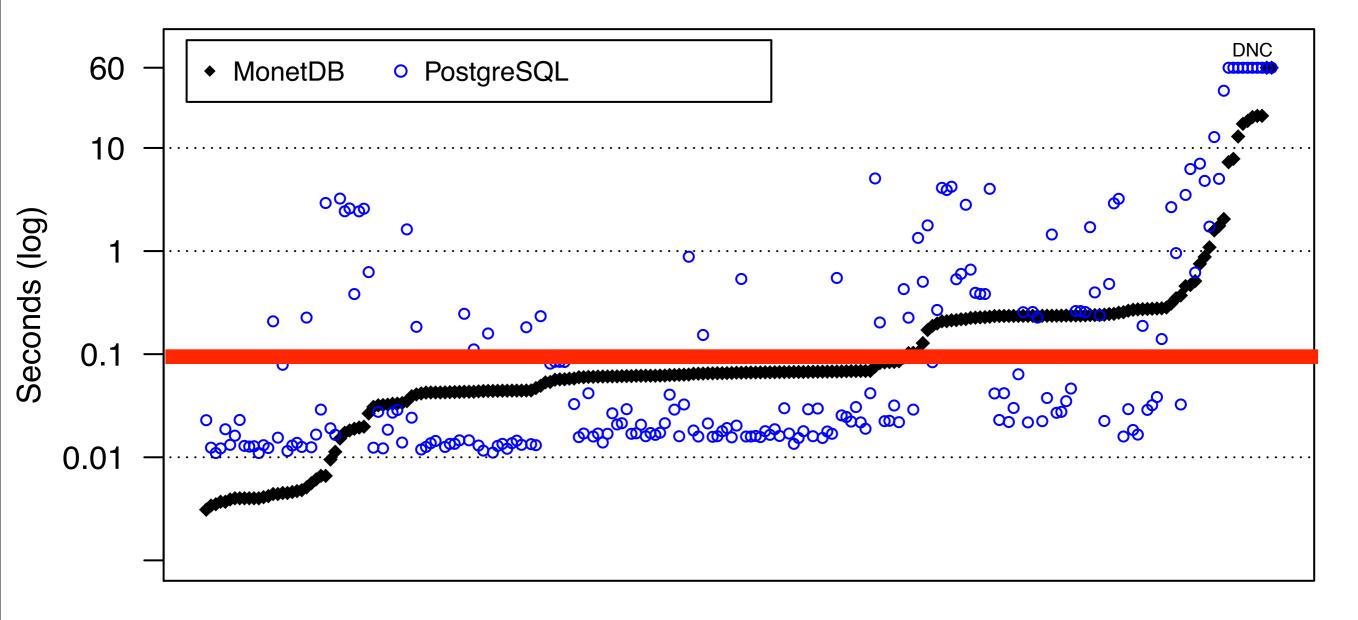


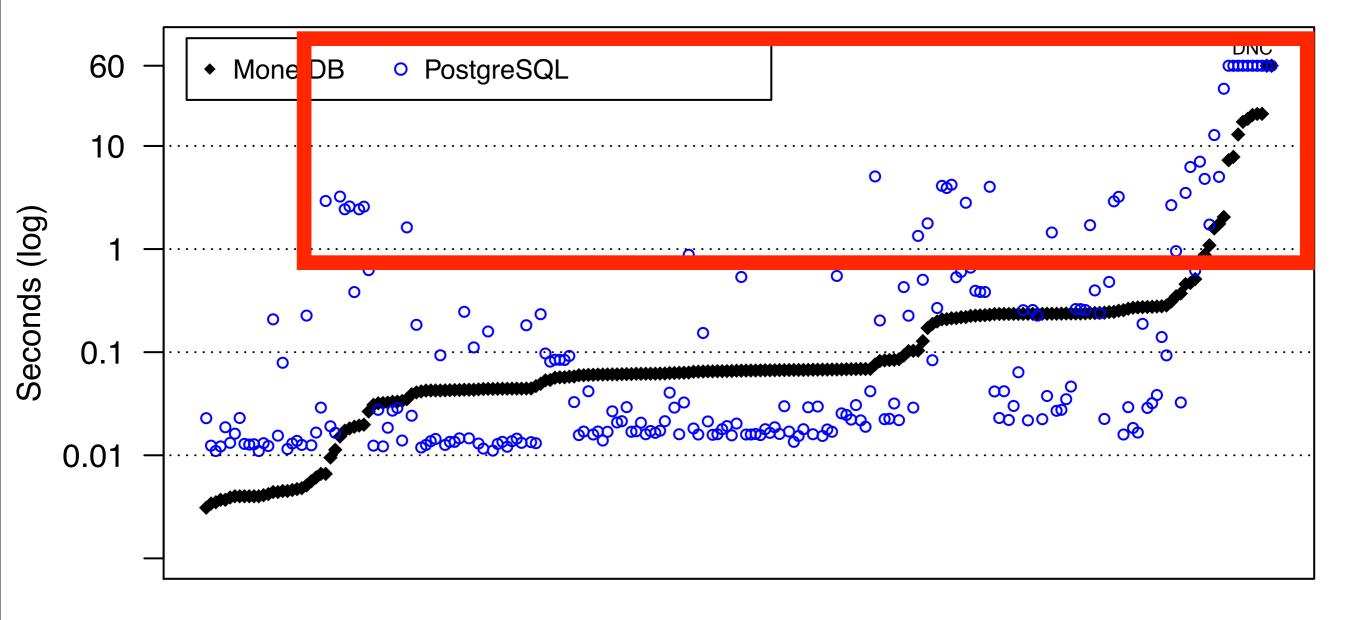
#### Workload of 10000 queries

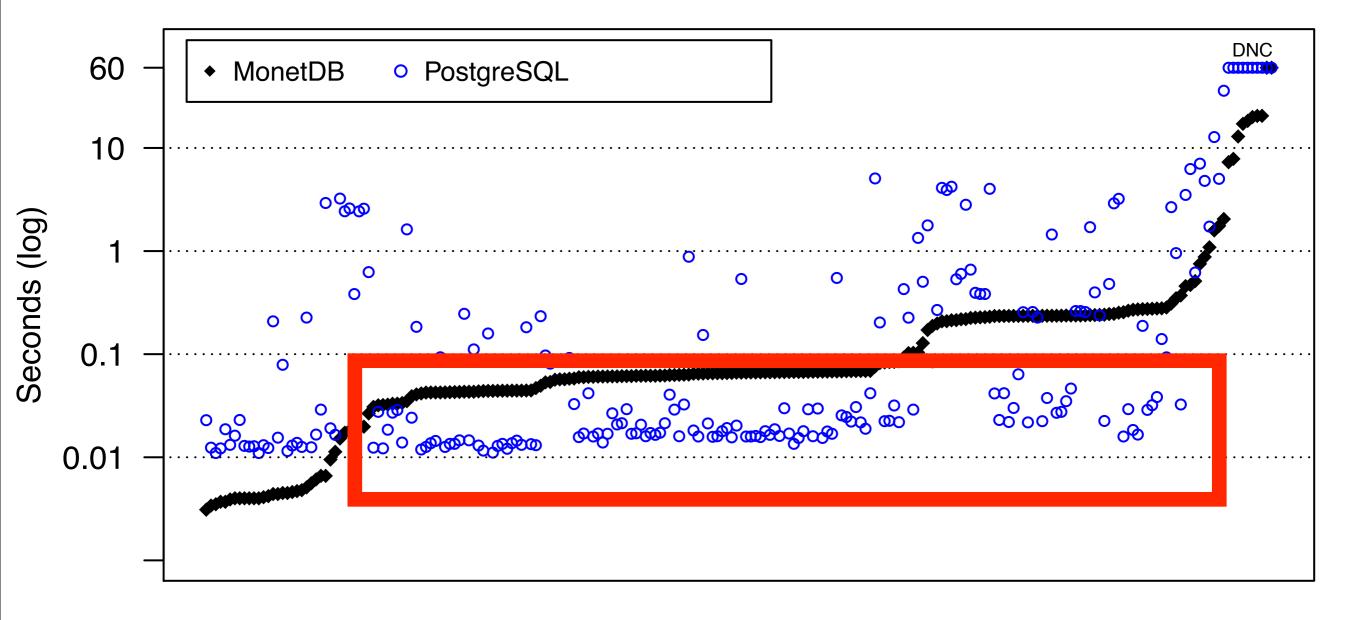




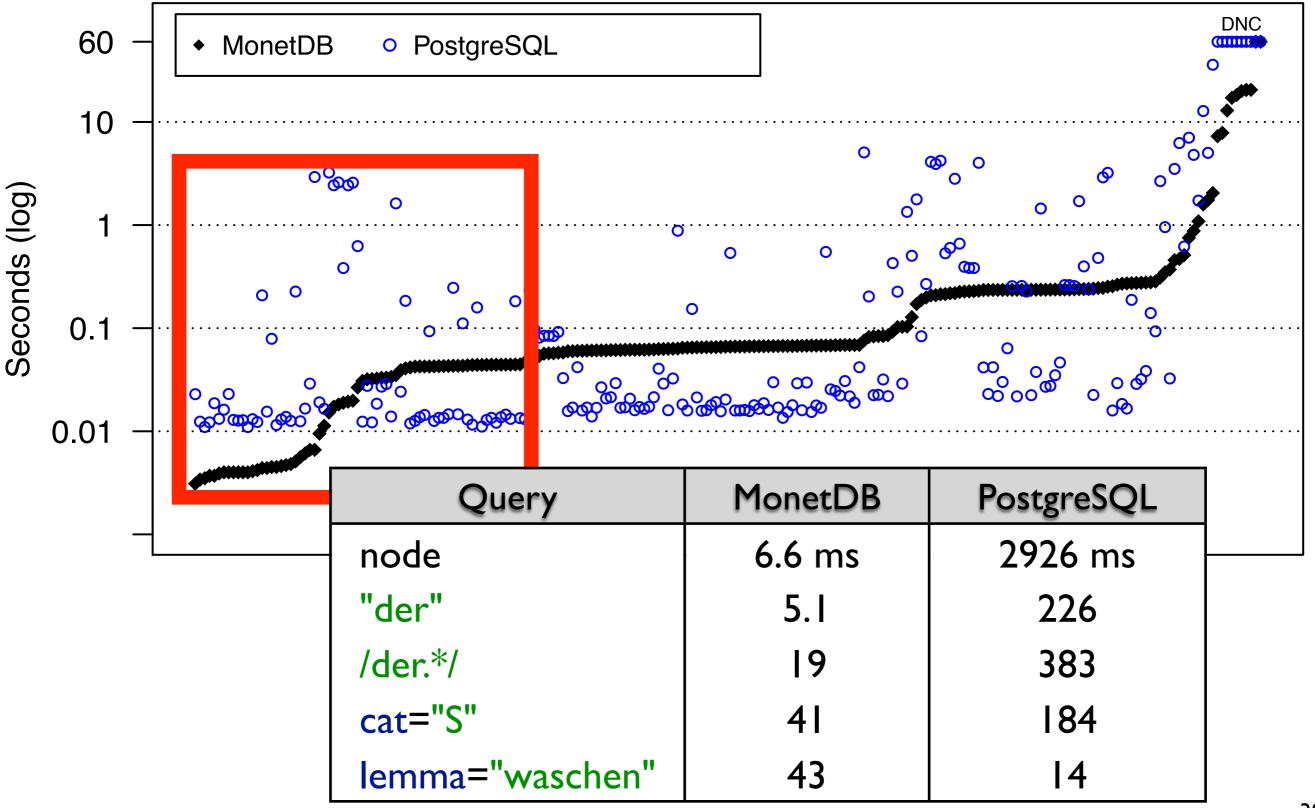




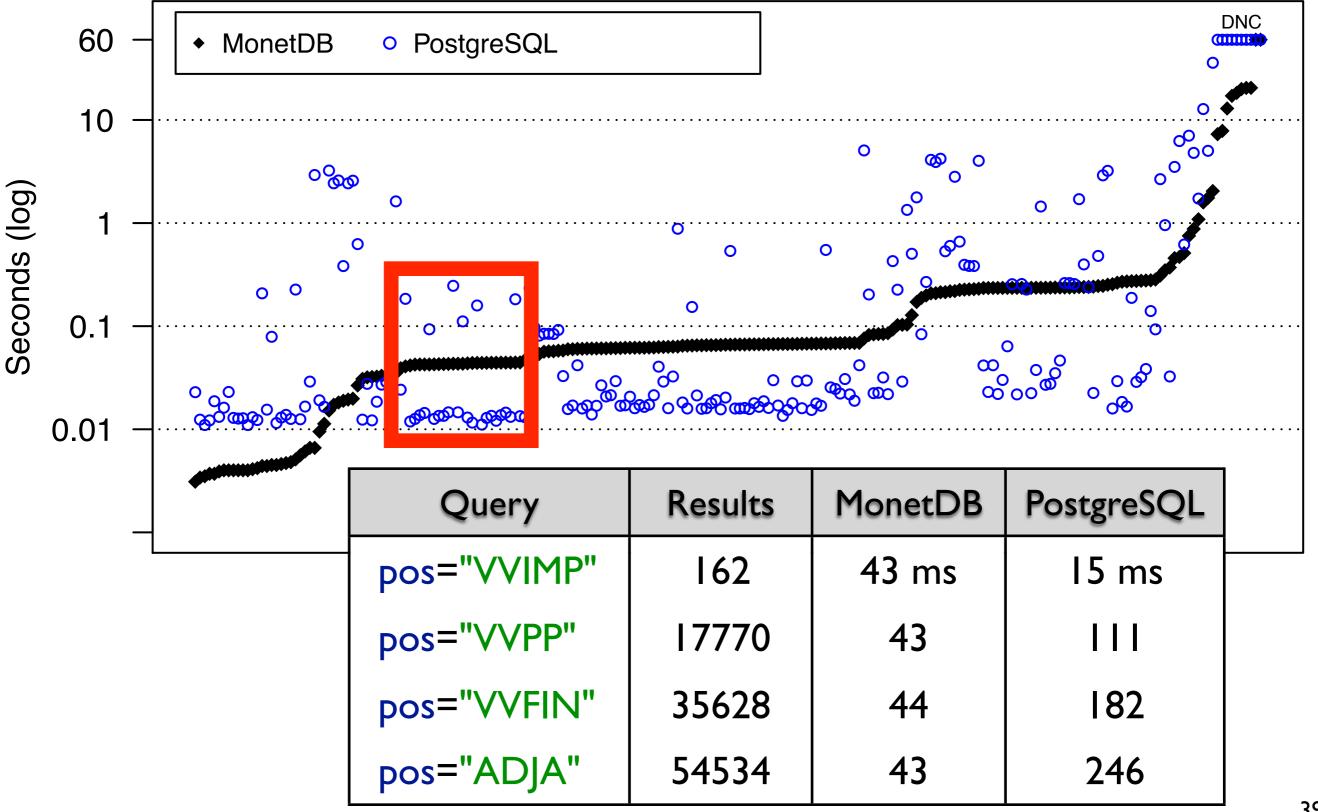




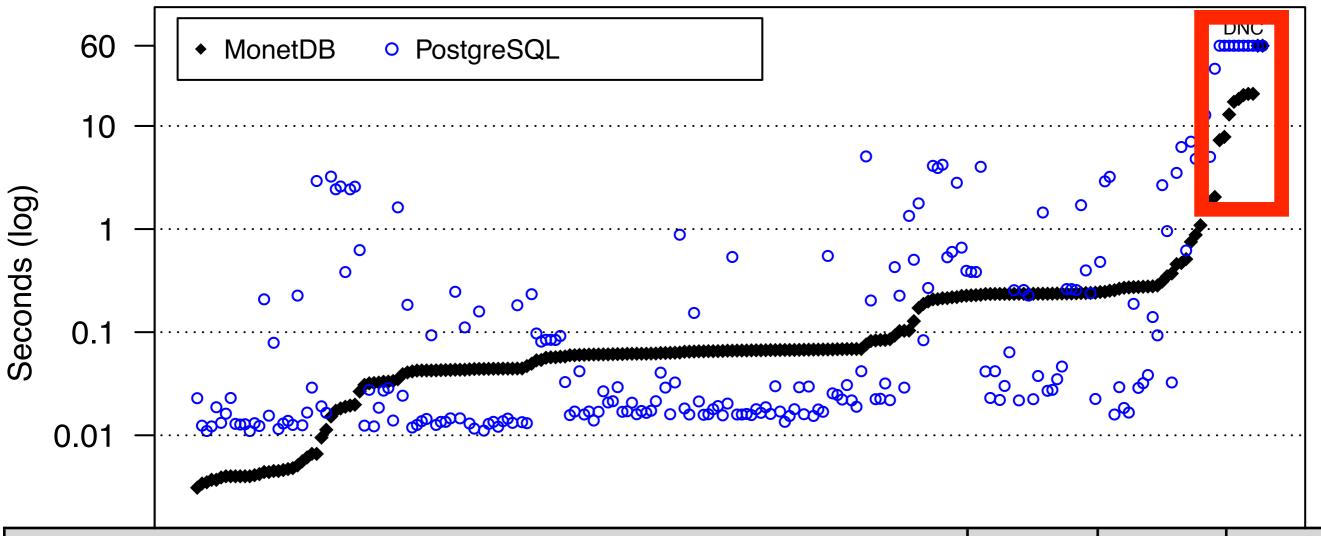
#### Simple queries are fast



#### Influence of result size



#### Queries with millions of results



Query		Monet	PSQL
lemma="müssen" & pos= /VV.*/ & pos="\$." & #1 .* #2 & #2 .* #3	4.5 M	2 s	35 s
pos=/VM.*/ & pos= /VV.*/ & pos=/.*/ & #1 .* #2 & #2 .* #3	384 M	175 s	> I h

#### Fast regular expressions

regular expression without a fixed prefix can't use an index, need to scan the entire column

Query	MonetDB	PostgreSQL
/.*sich.*/	213 ms	4206 ms
/[Kk]ann.*/	219	2812
pos="VVPP" & lemma=/(ge)?kommen/ & #I _=_ #2	229	383
pos=/N.*/ & /[12][09][0-9][0-9]/ & #  _=_ #2	246	2902
lemma=/[^äöü]+/ & /.+[äöü].+/ & pos="NN" & #  _=_ #2 & #2 _=_ #3	469	6246

#### Advantages

#### **MonetDB**

- better overall performance
- stable query performance
- fast regular expressions
- normalized schema
- greatly reduced disk consumption
- better use of limited resources

#### **PostgreSQL**

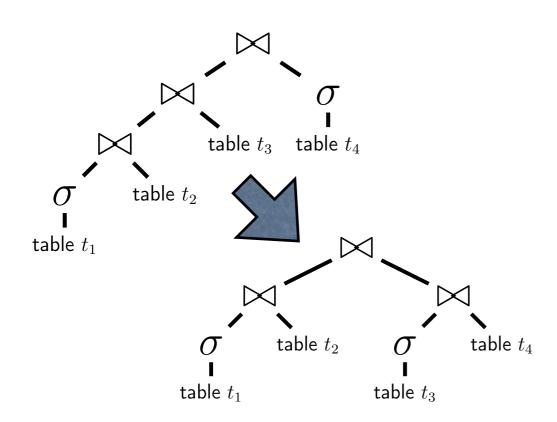
- queries with highly selective search term
- complete implementation
- bug-free SQL processing

#### Summary

- prototypical implementation of Annis on MonetDB
- test scenario from an Annis installation in service
- in-depth performance comparison of Annis on MonetDB and PostgreSQL

# SELECT vielen FROM dank;

#### Port of Annis to MonetDB

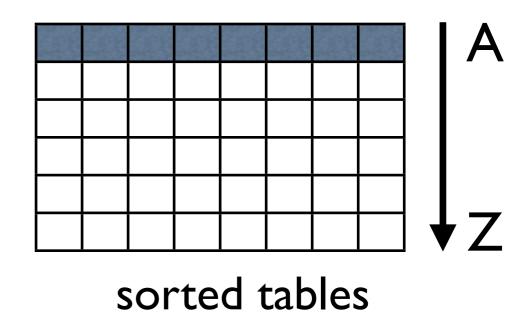


 $(a|b)^*$ 

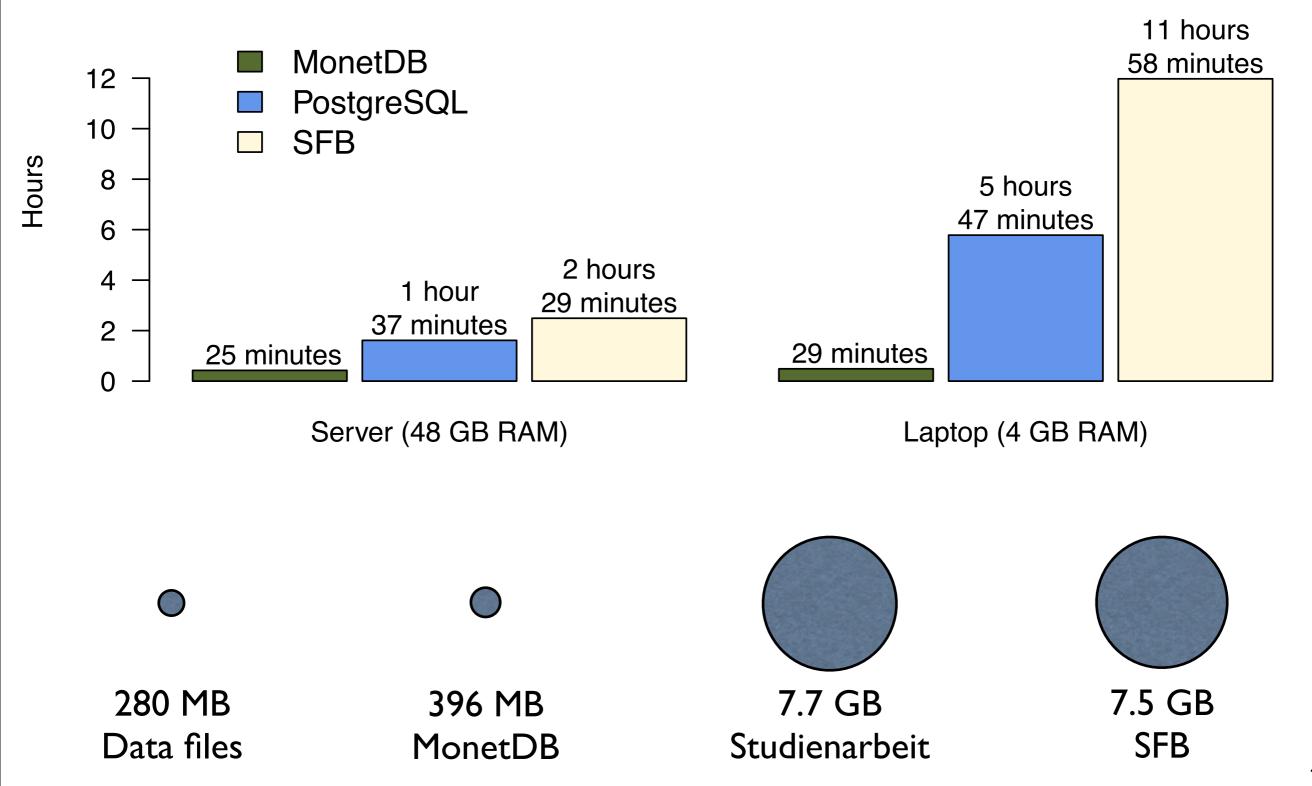
regular expressions

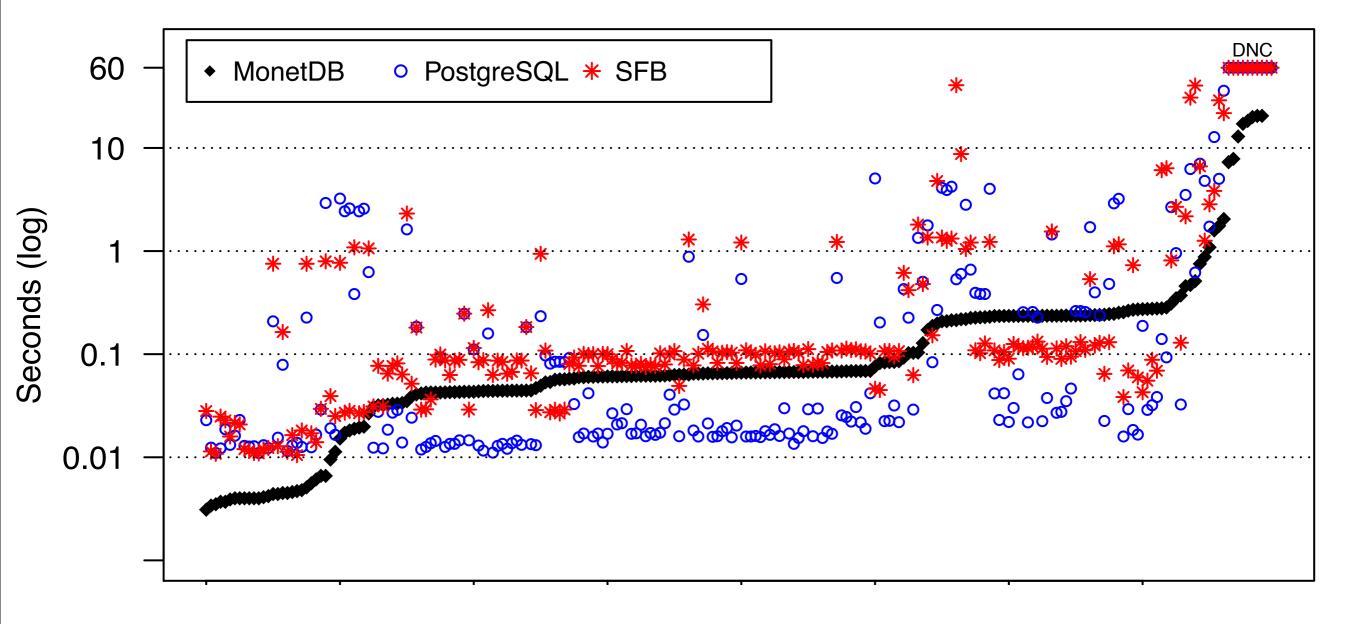
execution plan



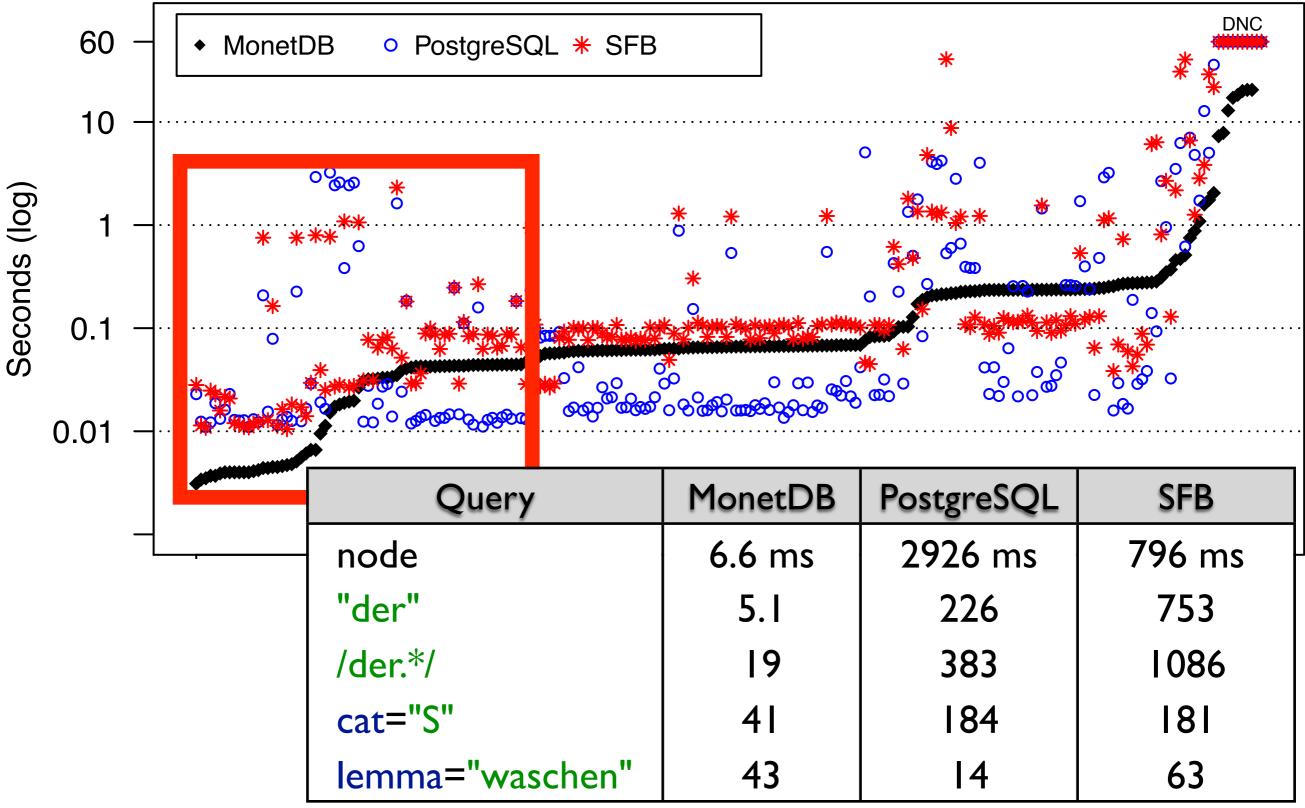


#### PostgreSQL vs. MonetDB

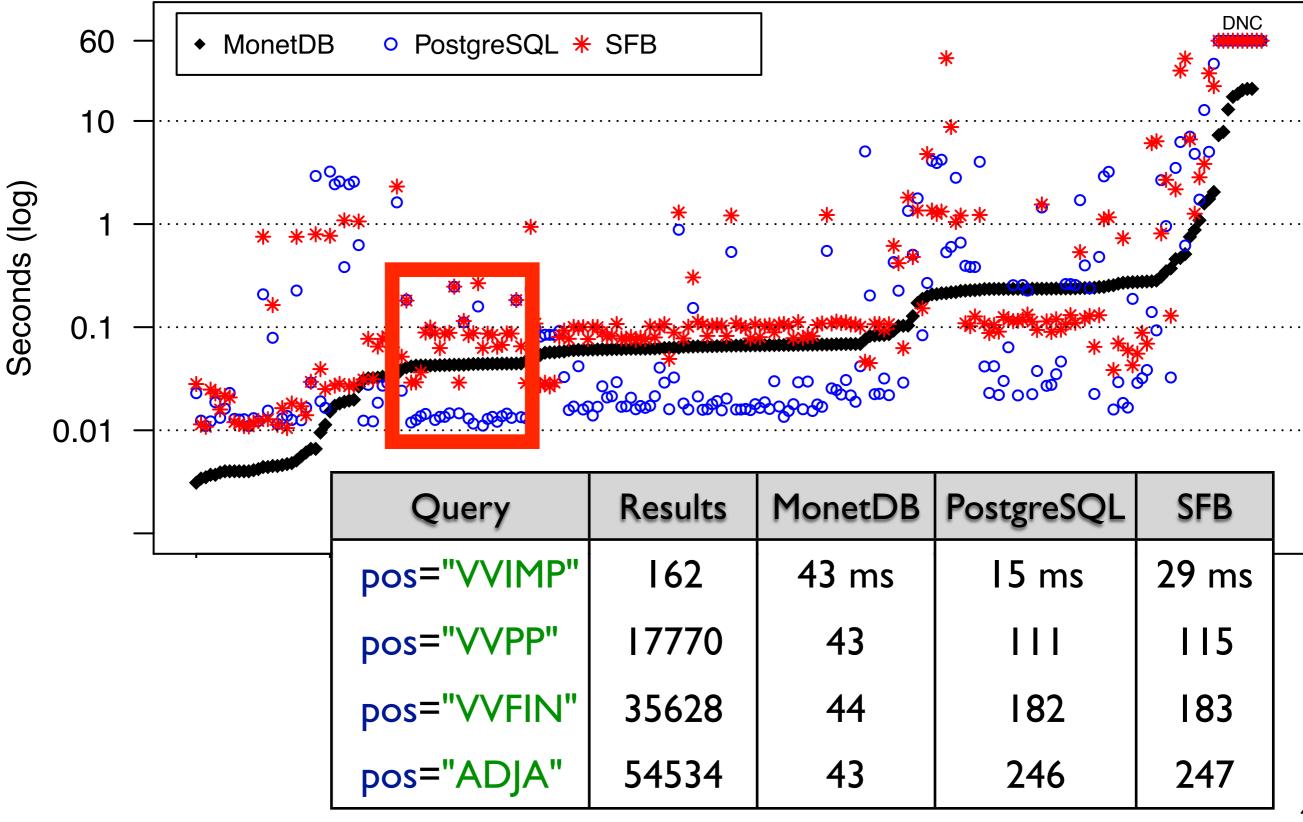




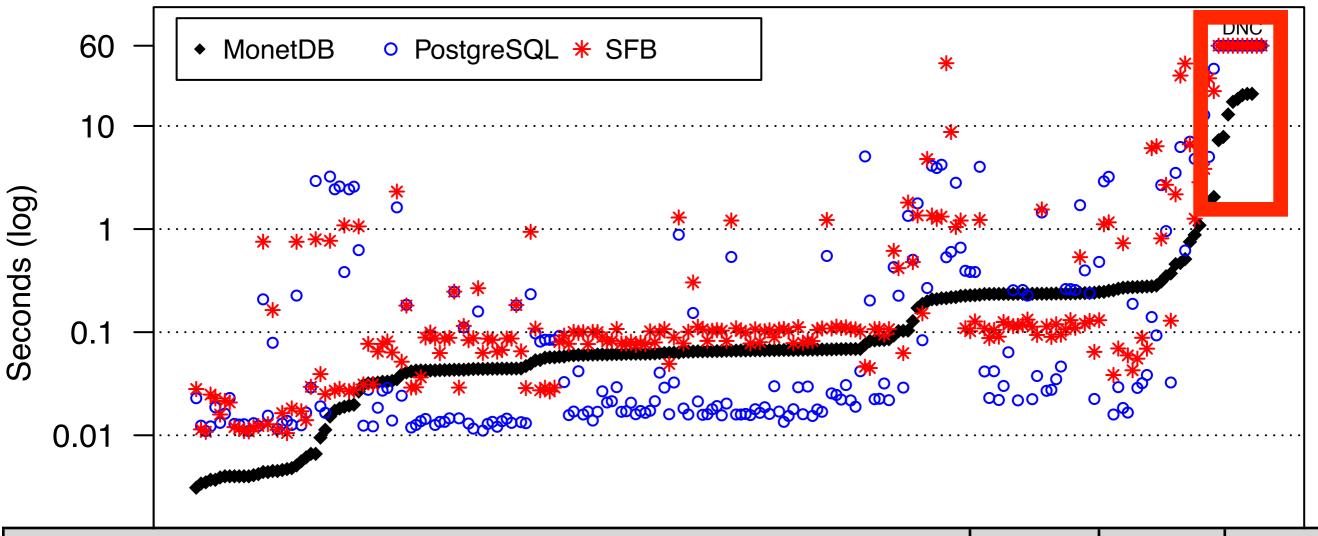
#### Simple queries



#### Influence of result size



#### Queries with millions of results



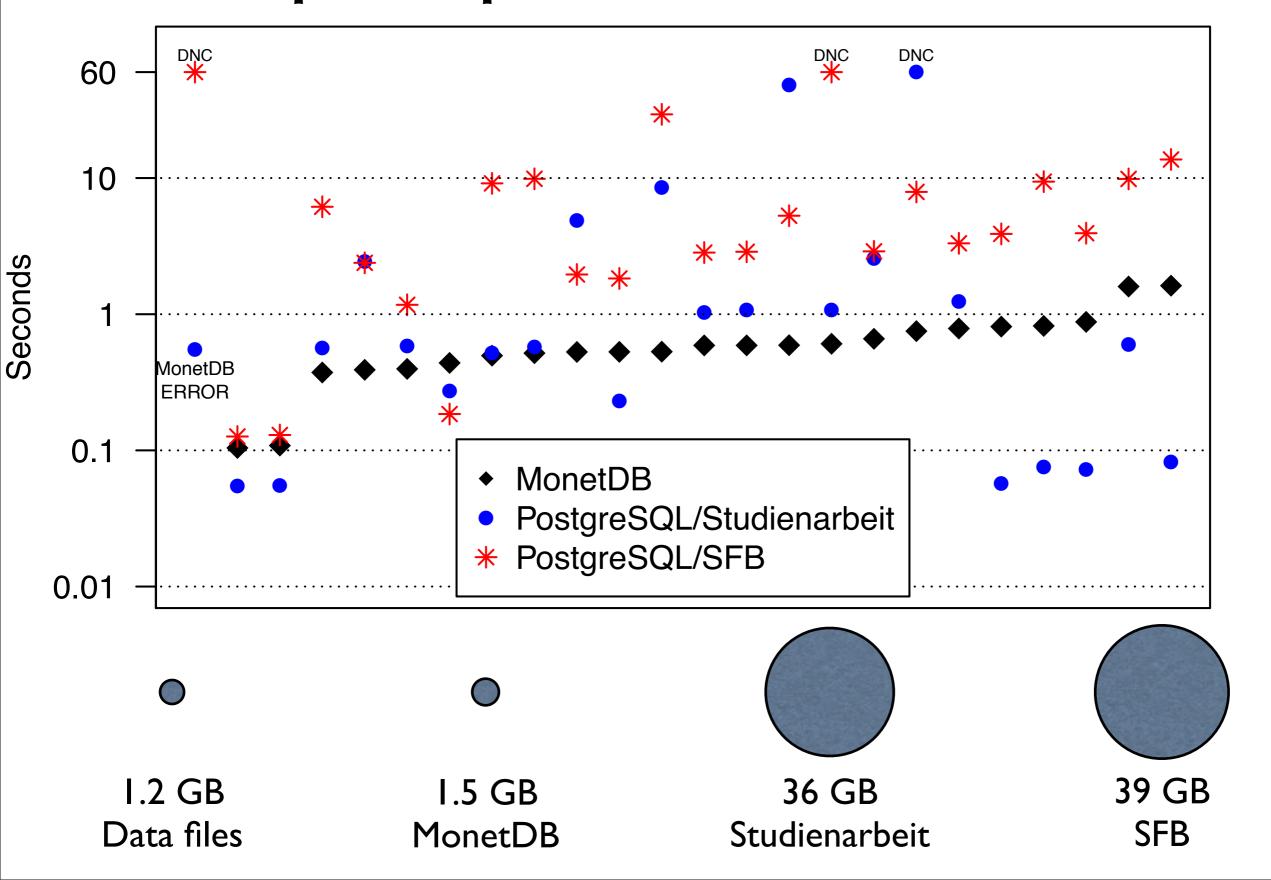
Query		PSQL	SFB
lemma="müssen" & pos= /VV.*/ & pos="\$." & #1 .* #2 & #2 .* #3	2 s	35 s	22 s
pos=/VM.*/ & pos= /VV.*/ & pos=/.*/ & #1 .* #2 & #2 .* #3	175 s	> I h	36 m

#### Regular expressions

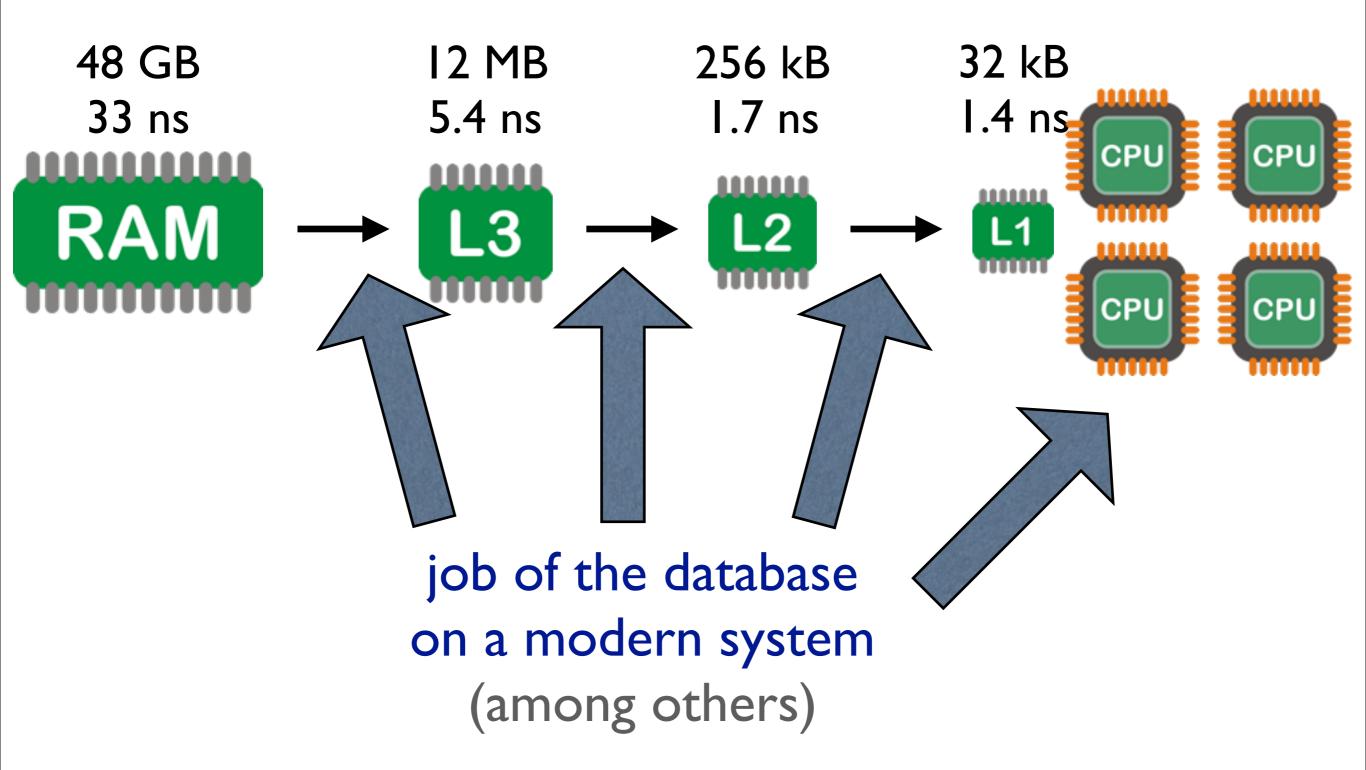
regular expression without a fixed prefix can't use an index, need to scan the entire column

Query	MonetDB	PSQL	SFB
/.*sich.*/	213 ms	4206 ms	1322 ms
/[Kk]ann.*/	219	2812	1047
pos="VVPP" & lemma=/(ge)?kommen/ & #I _=_ #2	229	383	127
pos=/N.*/ & /[12][09][0-9][0-9]/ & #I _=_ #2	246	2902	1111
lemma=/[^äöü]+/ & /.+[äöü].+/ & pos="NN" & #I _=_ #2 & #2 _=_ #3	469	6246	30803

#### Example queries on TüBa-D/Z

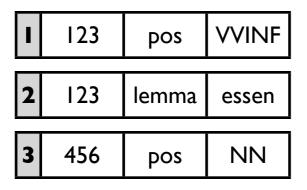


#### Multiple CPUs



# Algorithms

#### **Row-Store**



- function calls in a loop getNextTuple getAttribute('name') compareAttribute('lemma')
- pipeline flush
- instruction cache miss

#### Column-Store

node_ref	name	value
123	pos	VVINF
123	lemma	essen
456	pos	NN

- array lookup in a loop
  column[index] == 'lemma'
- loop unrolling
- no jumps
- instruction locality

#### 50% of time spent in resource stalls

#### History of Column-Stores

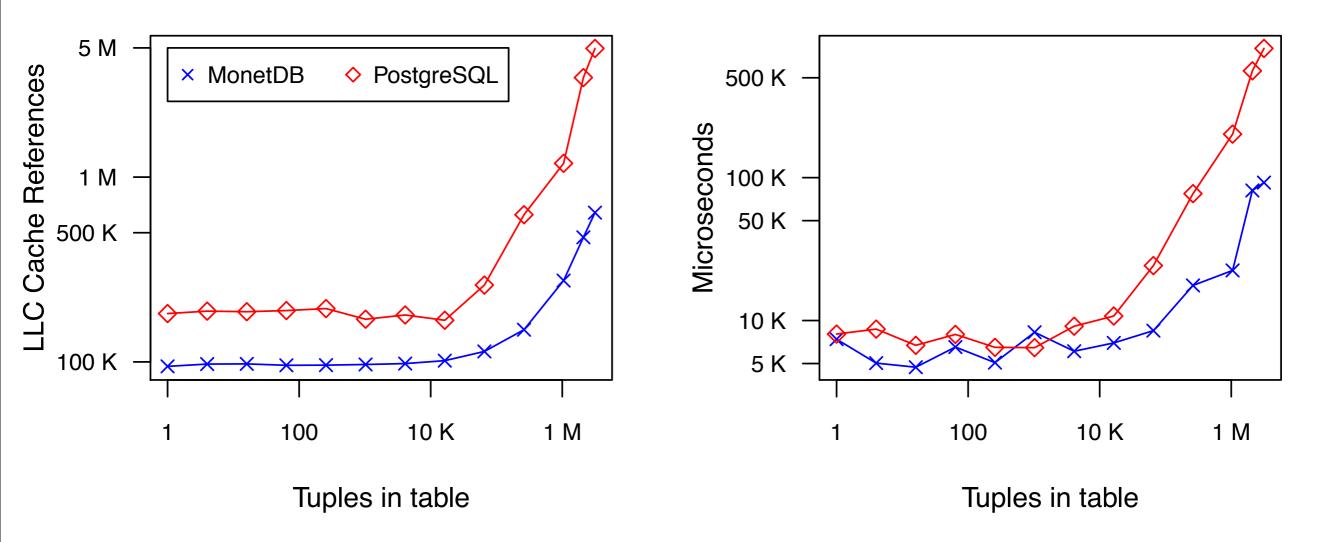
- IBM, 1973: System R
- Copeland, 1985: A Decomposed Storage Model
- CWI, 1994: Monet. An Impressionist Sketch of an Advanced Database system
- 2000s: Column-Stores become mainstream
  - C-Store, Vertica
  - MonetDB/X100, Vectorwise
  - SAP HANA

#### MonetDB

- pioneered research into Column-Stores
- active research platform
  - adaptive indexing, database cracking
  - partial vectorized bulk-processing, X100
  - opportunistic materialization, recycler
  - array primitives, SciQL
- well-received
  - VLDB 10 Years Best Paper Award, 2009: Database Architecture Optimized For The New Bottleneck: Memory Access
  - SIGMOD Jim Gray Doctoral Dissertation Award, 2011: Database Cracking: Towards Auto-tuning Database Kernels
  - VLDB Challenges & Vision Best Paper Award, 2011: The Researcher's Guide to the Data Deluge: Querying a Scientific Database in Just a Few Seconds



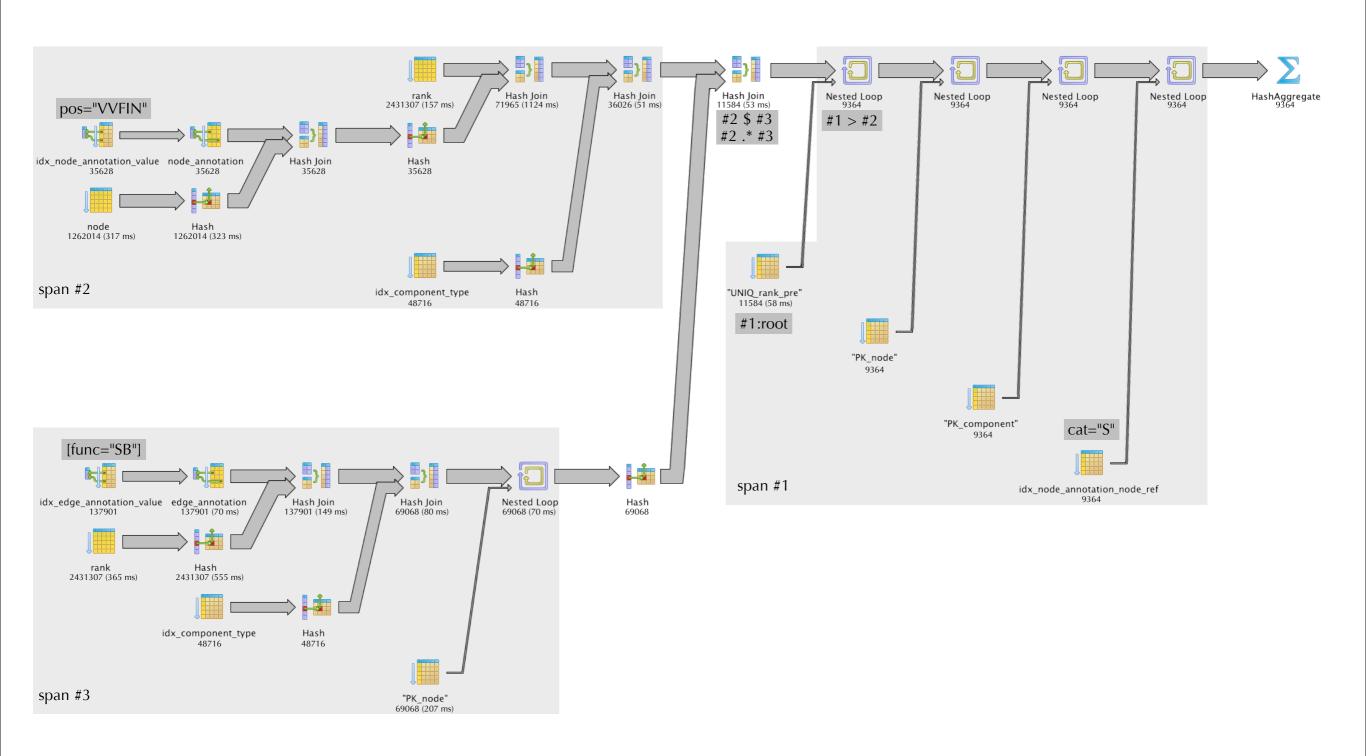
# Cache usage vs. query speed



Query: annotation name = 'lemma'

- last level cache references
- indirect measure of processed data
- everything that is processed must pass once through the cache

#### Normalized schema on PostgreSQL



#### Materialized schema on PostgreSQL

